

Solutions to HW-2

- Arjun Parasuram Prasad (NetID: ap9334)

Support Vector Machines: SVMs with Pegasos

Ans 1:

Ans 1 $g \in \mathbb{R}^d$ is a subgradient of $f: \mathbb{R}^d \rightarrow \mathbb{R}$
 if $\forall z$

$$f(z) \geq f(x) + g^T(z-x)$$

$f_1, f_2, f_3, \dots, f_m: \mathbb{R}^d \rightarrow \mathbb{R}$
 \hookrightarrow convex functions.

$$f(x) = \max_{i=1,2,\dots,m} f_i(x)$$

Let k be any index for which $f_k(x) = f(x)$
 and choose $g \in \partial f_k(x)$
 To prove: $g \in \partial f(x)$ also.

Proof:

given that: $g \in \partial (f_k(x))$
 & say \mathcal{D}_f is the domain of $f(x)$.

Let $\mathcal{D}_{f_{k_i}}$ be the domain of f_{k_i} which has subgradient g_{k_i} &

$$f_{k_p}(x) = \max_{i=1 \dots m} f_i(x)$$

Case 1: $z \in \mathcal{D}_{f_{k_0}}$ and $x \in \mathcal{D}_{f_{k_0}}$

\Rightarrow for subgradient g_{k_0} , we know that

$$f_{k_0}(z) \geq f_{k_0}(x) + g_{k_0}^T(z-x)$$

as $f_{k_0}(t) = f(t) \quad \forall t \in \mathcal{D}_{f_{k_0}}$

\therefore we can write this expression as

$$f(z) \geq f(x) + g^T(z-x)$$

\therefore when z and x belong to the same subdomain of f , \exists subgradient $g_{k_0} \in \partial(f(x))$

- Case 2. z and x belong to different subdomains of $f(x)$

Let $z \in \mathcal{D}_{f_{k_2}}$ and $x \in \mathcal{D}_{f_{k_1}}$.

Let $g_{k_1} \in \partial(f_{k_1}(x))$

Let $g_{k_2} \in \partial(f_{k_2}(z))$

Then for function f_{k_1} and subgrad. g_{k_1} we know that

$$f_{k_1}(z) \geq f_{k_1}(x) + g_{k_1}^T(z-x)$$

$$\text{as } z \in \mathcal{D}_{f_{k_2}} \Rightarrow f_{k_2}(z) \geq f_{k_1}(z)$$

∴

$$f_{k_2}(z) \geq f_{k_1}(x) + g_{k_1}^T(z-x)$$

$$\text{for } z \in \mathcal{D}_{f_{k_2}}, f_{k_2}(z) = f(z)$$

for $x \in D_{f_{k_1}}$, $f_{k_1}(x) = f(x)$

Thus, we can rewrite our eqⁿ as

$$f(z) \geq f(x) + g_{k_1}^T(z-x)$$

o.o even when z, x belong to
2 different subdomains of f ,

$\exists g_{k_1}$ such that $g_{k_1} \in \partial(f(x))$

Thus we have that

$\forall x, z \in D_f$
if $g_{k_1} \in \partial(f(x))$

Hence Proved.

Ans 2:

Ans 2

$$J(w) = \max(0, 1 - yw^T x)$$


\Rightarrow defining $g \in \partial(J(w))$

we can write

$$g = \begin{cases} -yx, & yw^T x < 1 \\ 0, & yw^T x \geq 1 \end{cases}$$

$[x \in \mathbb{R}^d,$
 $w \in \mathbb{R}^d,$
 $y \in \mathbb{R}]$

Proving that g is a subgradient for $J(w)$:-



$J(w)$ can be rewritten as

$$J(w) = \begin{cases} 1 - yw^T x, & yw^T x < 1 \\ 0, & yw^T x \geq 1 \end{cases}$$

Take points w_1 and $w_2 = w_1 + d$. [$w_1 \neq w_2$]

Case 1: $yw_1^T x \geq 1$ and $yw_2^T x \geq 1$ $\in \mathcal{D}_{J(w)}$
domain

$$\begin{aligned} \xrightarrow{\text{RHS}} J(w_1) + g^T(w_2 - w_1) \\ = 0 + 0 \geq 0 = J(w_2) \end{aligned}$$

$$\Rightarrow J(w_2) \geq J(w_1) + g^T(w_2 - w_1)$$

$\therefore g$ satisfies subgradient conditions for case 1

Case 2: $yw_1^T x \geq 1$ and $yw_2^T x < 1$

$$\begin{aligned} \text{RHS: } J(w_1) + g^T(w_2 - w_1) \\ = 0 + 0 = 0 \quad [g = 0 @ w_1] \end{aligned}$$

$$\text{LHS: } J(w_2) = 1 - yw_2^T x > 0$$

$$\therefore J(w_2) \geq J(w_1) + g^T(w_2 - w_1)$$

Thus g satisfies subgradient cond's of $J(w)$ for case 2.

Case 3:

$$y w_2^T x > 1 \text{ and } y w_1^T x < 1$$

RHS:

$$\Rightarrow J(w_1) + g^T(w_2 - w_1)$$

$$= 1 - y w_1^T x + (-y x)^T (w_2 - w_1)$$

$$= 1 - \cancel{y w_1^T x} - y w_2^T x + \cancel{y w_1^T x} \quad \left[\begin{array}{l} g = yx \text{ @ } w_1, \\ (-y x)^T w = -y w^T x \end{array} \right]$$

$$= 1 - y w_2^T x < 0$$

$$(\text{since } y w_2^T x > 1 \Rightarrow 1 - y w_2^T x < 0)$$

LHS:

$$J(w_2) = 0 > 1 - y w_2^T x$$

\Rightarrow

$$J(w_2) > J(w_1) + g^T(w_2 - w_1)$$

∴ g satisfies subgradient conditions for $J(w)$ for case 3.

$$\text{Case 4: } y w_1^T x < 1 \text{ and } y w_2^T x < 1$$

$$\Rightarrow \text{RHS} = J(w_1) + g^T(w_2 - w_1)$$

$$RHS = (1 - yw_1^T x) + (-yx)^T (w_2 - w_1)$$

$$\begin{aligned} [g &= -yx @ w_1] \\ [-y^T x^T w &= -y^T w^T x] \end{aligned}$$

\Rightarrow

$$RHS = 1 - \cancel{yw_1^T x} - yw_2^T x + \cancel{yw_1^T x}$$

$$= 1 - yw_2^T x = J(w_2) = LHS$$

\therefore

$J(w_2) \geq J(w_1) + g^T (w_2 - w_1)$
 hence g satisfies subgradient conditions
 for $J(w)$ for case 4.

Since g satisfies subgradient condition
 $J(z) = J(w) + g^T (z - w)$

$\forall z, w$ in the domain of $J(w)$

\therefore we have proved that

$$g \in \partial(J(w))$$

Ans 3:

* SVM with Pegasos algorithm.

Ans 3

$$J^p(w) = \frac{\lambda}{2} \|w\|^2 + \max\{0, 1 - y_i^* w^T x_i^*\}$$

The above expression is not defined for

$$(x_i, y_i) \in \{(x_i, y_i) \mid y_i^* w^T x_i^* = 1\}$$

$\forall w$, s.t. $y_i^* w^T x_i^* \neq 1$, we have

$$\nabla_w (J^p(w)) = \nabla_w \left(\frac{\lambda}{2} w^T w \right) + \nabla_w (\max\{0, 1 - y_i^* w^T x_i^*\})$$

$$= \begin{cases} \lambda w + (-y_i^* x_i^*), & 0 > 1 - y_i^* w^T x_i^* \\ \lambda w + 0 & 0 < 1 - y_i^* w^T x_i^* \end{cases}$$

$$= \begin{cases} \lambda w - y_i^* x_i^*, & 0 > 1 - y_i^* w^T x_i^* \\ \lambda w & 0 < 1 - y_i^* w^T x_i^* \end{cases}$$

please turn over

Ans 4:

Ans 4

$$J(w) = \frac{\lambda}{2} w^T w + \max \{0, 1 - y_i^* w^T x_i^*\}$$

$$\partial(J(w)) = \underbrace{\lambda w}_{\partial(f_1(w))} + \underbrace{\partial[\max \{0, 1 - y_i^* w^T x_i^*\}]}_{\partial(f_2(w))}$$

we know that if

$$f = f_1 + f_2 + f_3 + \dots + f_n$$

$$\text{then } \partial(f) = \partial(f_1) + \partial(f_2) + \dots + \partial(f_n)$$

$$\therefore \partial(J(w)) = \partial(f_1(w)) + \partial(f_2(w))$$

$$\nabla_w(f_1(w)) \in \partial(f_1(w))$$

[as $f_1(w)$ is convex and differentiable at all points]

\Rightarrow a subgradient $g_1(\omega)$ of $f_1(\omega)$ is

$$g_1(\omega) = \lambda \omega$$

we also know that a subgradient of $f_2(\omega)$ is

$$g_2 = \begin{cases} -x_i^* y_i^* & , y_i^* \omega^T x < 1 \\ 0 & , y_i^* \omega^T x \geq 1 \end{cases}$$

(refer proof in Ans 2 attached with this q. as well)

$$\therefore \partial(\mathcal{J}(\omega)) = \partial(f_1(\omega)) + \partial(f_2(\omega))$$

\Rightarrow a subgradient of $\mathcal{J}(\omega)$ is

$$g = \begin{cases} \lambda \omega - x_i^* y_i^* & , y_i^* \omega^T x < 1 \\ \lambda \omega & , y_i^* \omega^T x \geq 1 \end{cases} \quad // \text{ hence proved.}$$

Ans 5:

```
def gen_sparse_bag_of_words(list_of_words):  
    return Counter(list_of_words)
```

please turn over

Ans 6:

```
X = list(map(lambda review : gen_sparse_bag_of_words(review[0 : len(review)
- 1]), data))
y = list(map(lambda review_label : review_label[-1], data))

X_train = X[0:1500]
X_test = X[1500:]
y_train = y[0:1500]
y_test = y[1500:]
```

please turn over

Ans 7:

```
def pegasos_v1(X_train, y_train, lambda_reg, epochs, verbose = True,
tolerance = 0.001):
    W = {}
    n = len(X_train)
    err = np.inf

    for t in range(1, n*epochs+1):
        if(t%epochs == 0):
            # reshuffle the data
            data = list(zip(X_train, y_train))
            random.shuffle(data)
            X_train, y_train = zip(*data)

            eta = 1/(lambda_reg*t)
            Xj = X_train[(t-1)%n]
            yj = y_train[(t-1)%n]

            margin = yj * dotProduct(W, Xj)

            if margin >= 1:
                for i, v in Xj.items():
                    W[i] = (1 - (eta*lambda_reg))*W.get(i, 0)
            else:
                for i, v in Xj.items():
                    W[i] = (1 - (eta*lambda_reg))*W.get(i, 0) + v*eta*yj

            if (t%n == 0) and (verbose == True):
                clf_error = classification_error(X_train, y_train, W)
                print('-----')
                print(f'epoch: {t/n}')
                print(f'classification error is: {clf_error}')
                print(f'W.size is : {len(W)}')
                if abs(err - clf_error) <= tolerance:
                    break
                err = clf_error

    return W
```

please turn over

Ans 8:

```
def pegasos_v2(X_train, y_train, lambda_reg, epochs, verbose = True,
tolerance = 0.001):
    W = {}
    s = 1
    n = len(X_train)
    err = np.inf

    for t in range(2, epochs*n + 1):
        if(t%epochs == 0):
            # reshuffle the data
            data = list(zip(X_train, y_train))
            random.shuffle(data)
            X_train, y_train = zip(*data)

            eta = 1 / (lambda_reg*t)
            Xj = X_train[(t-1)%n]
            yj = y_train[(t-1)%n]
            margin = yj * dotProduct(W, Xj) * s

            s = (1 - eta*lambda_reg)*s
            if margin < 1:
                increment(W, (1/s)*eta*yj, Xj)

            if(t%n == 0 and verbose == True):
                clf_error = classification_error(X_train, y_train, scale(W, s))
                print('-----')
                print(f'epoch: {t/n}')
                print(f'classification error is: {clf_error}')
                print(f'W.size is : {len(W)}')
                if abs(err - clf_error) <= tolerance:
                    break
                err = clf_error
    return scale(W, s)
```

please turn over

Ans 9:

```
print('epochs = 2')
begin = time.time()
pegasos_v1(X_train, y_train, lambda_reg= 0.01, epochs= 2, verbose= False,
tolerance= 0.001)
print(f'training time - pegasos_v1: {time.time() - begin}')
begin = time.time()
pegasos_v2(X_train, y_train, lambda_reg= 0.01, epochs= 2, verbose= False,
tolerance= 0.001)
print(f'training time - pegasos_v2: {time.time() - begin}')

print('epochs = 200. Tolerance has been disabled')
begin = time.time()
pegasos_v1(X_train, y_train, lambda_reg= 0.01, epochs= 200, verbose= False,
tolerance= -1)
print(f'training time - pegasos_v1: {time.time() - begin}')
begin = time.time()
pegasos_v2(X_train, y_train, lambda_reg= 0.01, epochs= 200, verbose= False,
tolerance= -1)
print(f'training time - pegasos_v2: {time.time() - begin}')
```

```
A9
epochs = 2
training time - pegasos_v1: 1.4444091320037842
training time - pegasos_v2: 1.3131685256958008
epochs = 200. Tolerance has been disabled
training time - pegasos_v1: 29.29502511024475
training time - pegasos_v2: 15.590850353240967
```

please turn over

Ans 10:

```
def classification_error(X, y, W):  
    y_pred = get_predictions(X, W)  
    return sum(np.array(y) != np.array(y_pred))/(len(y_pred))
```

```
print(f"validation set classification error using W_pegasos_v2: \  
      {classification_error(X_test, y_test, W_pegasos_v2)}")
```


```
validation set classification error using W_pegasos_v2: 0.178
```

please turn over

Ans 11:

```
def plot_lambda_vs_error(X_train, y_train, X_test, y_test, lambda_reg_list
= [], epochs = 100):
    error_list = []
    for lambda_reg in lambda_reg_list:
        W = pegasos_v2(X_train, y_train, lambda_reg= lambda_reg, epochs=
epochs, tolerance = -1, verbose=False)
        error_list.append(classification_error(X_test, y_test, W)*100)
    plt.figure(figsize=(10, 10))
    plt.plot([np.log10(i) for i in lambda_reg_list], error_list, 'bo-',
label='Lambda vs %Classification Error')
    plt.xlabel("log_lambda_reg")
    plt.ylabel("% classification error")
    plt.show()
```

```
plot_lambda_vs_error(X_train, y_train, X_test, y_test, \
    [1e-10, 1e-9, 1e-7, 1e-6, 1e-5, 1e-4, 1e-3, 1e-2, 1e-1, 1, 10, 100],
epochs=100)
```



please turn over


Ans 12:

```
def plot_score_grps_vs_percentage_error(X, y, W, bins):
    scores = [dotProduct(i, W) for i in X]
    coupled_mat = np.column_stack((y, scores))
    coupled_mat = coupled_mat[coupled_mat[:, 1].argsort()[::-1]]
    range_start = 0
    range_end = 0
    score_ranges = []
    classification_errors = []

    for i in np.linspace(0, coupled_mat.shape[0], bins, dtype=int):
        if i == range_end:
            continue
        range_start = range_end
        range_end = i
        mat = coupled_mat[range_start:range_end, :]
        score_ranges.append(f"{np.min(mat[:, 1]):.2f} to {np.max(mat[:, 1]):.2f}")
        classification_errors.append(100 * (np.sum(mat[:, 0] !=
np.sign(mat[:, 1]))) / mat.shape[0])

    plt.figure(figsize=(10, 10))
    plt.plot(score_ranges, classification_errors, 'ro-', label='score
ranges vs classification erros')
    plt.xlabel("score ranges")
    plt.ylabel("%classification error")
    plt.xticks(fontsize=7)
    plt.yticks(fontsize=7)
    plt.show()
```

```
W_pegasos_v2 = pegasos_v2(X_train, y_train, lambda_reg= 1e-4, epochs= 50,
verbose= False, tolerance= 1e-3)
plot_score_grps_vs_percentage_error(X_test, y_test, W_pegasos_v2, bins= 8)
```



Observations

- Classification error is relatively low for those score ranges which have high magnitude, i.e. $\text{abs}(\text{score})$ is high.
- Classification error is relatively high for those score ranges which have lower magnitude of scores, i.e. $\text{abs}(\text{score})$ is low

please turn over

Kernel Methods

Ans 13:

```

### Kernel function generators
def linear_kernel(X1, X2):
    """
    Computes the linear kernel between two sets of vectors.
    Args:
        X1 - an n1xd matrix with vectors x1_1,...,x1_n1 in the rows
        X2 - an n2xd matrix with vectors x2_1,...,x2_n2 in the rows
    Returns:
        matrix of size n1xn2, with  $x1_i^T x2_j$  in position i,j
    """
    return np.dot(X1,np.transpose(X2))

def RBF_kernel(X1,X2,sigma):
    """
    Computes the RBF kernel between two sets of vectors
    Args:
        X1 - an n1xd matrix with vectors x1_1,...,x1_n1 in the rows
        X2 - an n2xd matrix with vectors x2_1,...,x2_n2 in the rows
        sigma - the bandwidth (i.e. standard deviation) for the
    RBF/Gaussian kernel
    Returns:
        matrix of size n1xn2, with  $\exp(-||x1_i-x2_j||^2/(2 \sigma^2))$  in
    position i,j
    """
    sq_euclidean_dist = scipy.spatial.distance.cdist(X1, X2, 'sqeuclidean')
    return np.exp((-1)*(1/(2*sigma*sigma))*sq_euclidean_dist)

def polynomial_kernel(X1, X2, offset, degree):
    """
    Computes the inhomogeneous polynomial kernel between two sets of
    vectors
    Args:
        X1 - an n1xd matrix with vectors x1_1,...,x1_n1 in the rows
        X2 - an n2xd matrix with vectors x2_1,...,x2_n2 in the rows
        offset, degree - two parameters for the kernel
    Returns:
        matrix of size n1xn2, with  $(\text{offset} + \langle x1_i, x2_j \rangle)^{\text{degree}}$  in
    position i,j
    """
    return (np.dot(X1,np.transpose(X2)) + offset)**degree

```

please turn over

Ans 14:

```
linear_kernel(np.array([[-4], [-1], [0], [2]]), np.array([[-4], [-1], [0], [2]]))
```

```
array([[16,  4,  0, -8],  
       [ 4,  1,  0, -2],  
       [ 0,  0,  0,  0],  
       [-8, -2,  0,  4]])
```

please turn over

Ans 15:

```

plot_step = .01
xpts = np.arange(-6.0, 6, plot_step).reshape(-1,1)
prototypes = np.array([-4,-1,0,2]).reshape(-1,1)

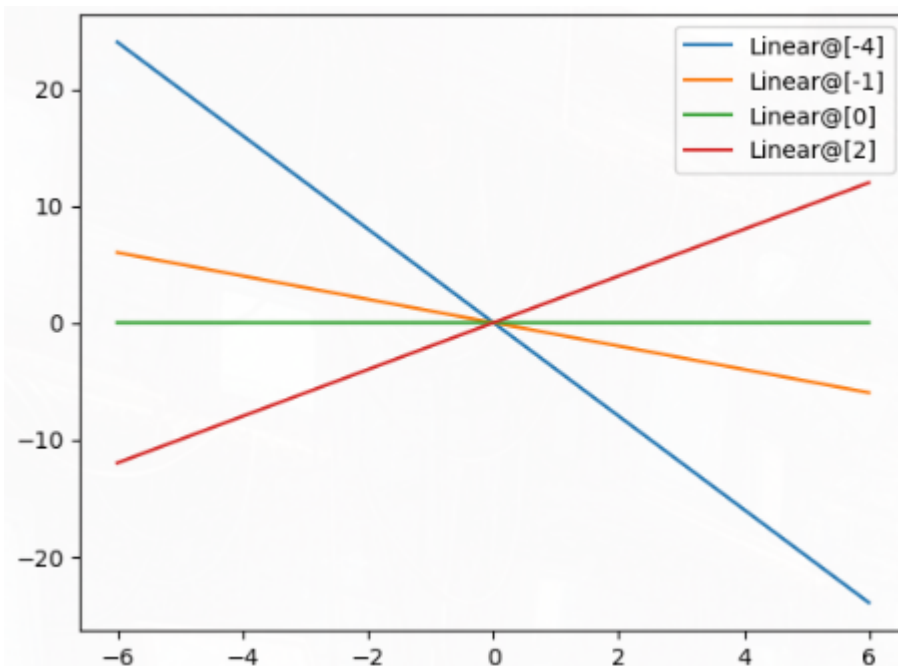
# Linear kernel
y = linear_kernel(prototypes, xpts)
for i in range(len(prototypes)):
    label = "Linear@"+str(prototypes[i,:])
    plt.plot(xpts, y[i,:], label=label)
plt.legend(loc = 'best')
plt.show()

# Polynomial kernel
y = polynomial_kernel(prototypes, xpts, offset=1, degree=3)
for i in range(len(prototypes)):
    label = "Polynomial@"+str(prototypes[i,:])
    plt.plot(xpts, y[i,:], label=label)
plt.legend(loc = 'best')
plt.show()


# RBF kernel
y = RBF_kernel(prototypes, xpts, sigma= 1)
for i in range(len(prototypes)):
    label = "RBF@"+str(prototypes[i,:])
    plt.plot(xpts, y[i,:], label=label)
plt.legend(loc = 'best')
plt.show()

```


- set of functions $x \rightarrow k_{\text{linear}}(x_0, x)$ for $x_0 \in DX$ and for $x \in [-6, 6]$



- set of functions $x \rightarrow \text{kpoly}(1,3)(x_0,x)$ for $x_0 \in DX$ and for $x \in [-6,6]$



- set of functions $x \rightarrow \text{kRBF}(1)(x_0,x)$ for $x_0 \in DX$ and for $x \in [-6,6]$



please turn over

Ans 16:

```

class Kernel_Machine(object):
    def __init__(self, kernel, training_points, weights):
        """
        Args:
            kernel(X1,X2) - a function return the cross-kernel matrix
            between rows of X1 and rows of X2 for kernel k
            training_points - an nxd matrix with rows x_1,..., x_n
            weights - a vector of length n with entries alpha_1,...,alpha_n
        """

        self.kernel = kernel
        self.training_points = training_points
        self.weights = weights

    def predict(self, X):
        """
        Evaluates the kernel machine on the points given by the rows of X
        Args:
            X - an nxd matrix with inputs x_1,...,x_n in the rows
        Returns:
            Vector of kernel machine evaluations on the n points in X.
            Specifically, jth entry of return vector is
            Sum_{i=1}^R alpha_i k(x_j, mu_i)
        """
        K = self.kernel(X, self.training_points)
        return np.dot(K, self.weights)

```


```

prototypes = np.array([-1, 0, 1]).reshape(-1,1)
weights = np.array([1, -1, 1]).reshape(-1, 1)

Fx = Kernel_Machine(kernel= (lambda i, j : RBF_kernel(X1 = i, X2 = j,
sigma=1)), \
    training_points= prototypes, weights=weights)

plot_step = .01
xpts = np.arange(-6.0, 6, plot_step).reshape(-1,1)
plt.plot(xpts, Fx.predict(xpts))
plt.show()

```



please turn over

Logistic Regression

Ans 17:

Ans 17

To prove: Show that ERM with logistic loss and MLE with a Bernoulli response distribution and the logistic link function will produce the same solution for w .

ERM with logistic loss

$$x \in \mathbb{R}^d, \quad w \in \mathbb{R}^d, \quad y_{\pm} = \{+1, -1\}$$

$$D = \{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}) \dots (x^{(n)}, y^{(n)})\}$$

$$\mathcal{F}_{\text{score}} = \{x \mapsto x^T w \mid w \in \mathbb{R}^d\}$$

$$\hat{y} = \text{sign}\{x^T w\}$$

logistic loss function

$$l_{\text{logistic}}(y, w) = \log(1 + \exp(-y w^T x))$$

$$\begin{aligned} \Rightarrow \hat{R}_n &= \frac{1}{n} \sum_{i=1}^n l_{\text{logistic}}(f(x^{(i)}), y^{(i)}) \\ &= \frac{1}{n} \sum_{i=1}^n \log(1 + \exp(-y^{(i)} w^T x^{(i)})) \end{aligned}$$

$$\begin{aligned} \Rightarrow \hat{R}_n &= \frac{1}{2n} \sum_{i=1}^n \left[(1+y^{(i)}) \log(1 + \exp(-w^T x^{(i)})) \right. \\ &\quad \left. + (1-y^{(i)}) \log(1 + \exp(-w^T x^{(i)})) \right] \end{aligned}$$

minimizing $\hat{R}_n \Rightarrow$

$$\hat{f} \in \arg \min_f R_n(f)$$

$$\Rightarrow \hat{w} \in \arg \min_{w \in \mathbb{R}^d} R_n(w)$$

$$\Rightarrow \hat{w} \in \arg \min_{w \in \mathbb{R}^d} \frac{1}{n} \sum_{i=1}^n \log(1 + \exp(-\underline{\underline{y^{(i)}} w^T x^{(i)}}))$$

\Downarrow (proof done in assignment ① as well. \rightarrow margin)

$$\Rightarrow \hat{w} \in \arg \min_{w \in \mathbb{R}^d} \left(\frac{1}{n} \sum_{i=1}^n \left[\frac{1}{2} (1 - y^{(i)}) \log(1 + \exp(-w^T x^{(i)})) + \frac{1}{2} (1 + y^{(i)}) \log(1 + \exp(w^T x^{(i)})) \right] \right)$$

eqn ①

MLE with a Bernoulli response distribution and the logistic link function.

$$\left[\begin{array}{l} x \in \mathbb{R}^d, y \in \{-1, 1\}, w \in \mathbb{R}^d \\ D = \{(x^{(1)}, y^{(1)}) (x^{(2)}, y^{(2)}) \dots (x^{(n)}, y^{(n)})\} \end{array} \right]$$

$$p(y=1 | x; w) = \frac{1}{1 + \exp(-x^T w)}$$

$$\begin{aligned} p(y=-1 | x; w) &= 1 - p(y=1 | x; w) \\ &= 1 - \frac{1}{1 + \exp(-x^T w)} \end{aligned}$$

$$\Rightarrow p(y=-1 | x; w) = \frac{1}{1 + \exp(x^T w)}$$

Likelihood function.

$$\Rightarrow L(\omega) = \prod_{i=1}^n p(y^{(i)} | x^{(i)}; \omega)$$

$$= \prod_{i=1}^n \left[\left(\frac{1}{1 + \exp(-x^{(i)\top} \omega)} \right)^{\left(\frac{1+y^{(i)}}{2} \right)} \left(\frac{1}{1 + \exp(x^{(i)\top} \omega)} \right)^{\left(\frac{1-y^{(i)}}{2} \right)} \right]$$

\Rightarrow the log likelihood function is :

$$\begin{aligned} \log L(\omega) = \frac{1}{2} \sum_{i=1}^n & \left[(1+y^{(i)}) \log((1 + \exp(-x^{(i)\top} \omega))^{-1}) \right. \\ & + \\ & \left. (1-y^{(i)}) \log((1 + \exp(x^{(i)\top} \omega))^{-1}) \right] \end{aligned}$$

[NOTE:

$$x^T w = w^T x \quad \text{as } x \in \mathbb{R}^d, w \in \mathbb{R}^d$$

]

The negative log likelihood can be written as :

$$NLL(w) = -\log l(w)$$

$$= \frac{1}{n} \sum_{i=1}^n \left[(1-y^{(i)}) \log(1 + \exp(-w^T x^{(i)})) + (1+y^{(i)}) \log(1 + \exp(w^T x^{(i)})) \right]$$

• minimizing the $NLL(w)$

$$\Rightarrow \hat{w} \in \arg \min_{w \in \mathbb{R}^d} NLL(w)$$

$$\Rightarrow \hat{w} \in \arg \min_{w \in \mathbb{R}^d} \left(\frac{1}{n} \right) NLL(w)$$

$$\Rightarrow \hat{w} \in \arg \min_{w \in \mathbb{R}^d} \left(\frac{1}{n} \left(\sum_{i=1}^n \left[(1-y^{(i)}) \log(1+\exp(-w^T x_i)) + (1+y^{(i)}) \log(1+\exp(w^T x_i)) \right] \right) \right)$$

eqn (2)

→ This is the same as eqn (1)

∴ since the optimization problem is same in both eqn.s (1) and (2)

∴ \hat{w} produced as a solution of eqn (1) and (2) will also be the same.

Hence proved.

Ans 18:

Ans 18

Decision boundary of logistic regression

$$= \log \frac{P(y=1|x;\omega)}{P(y=-1|x;\omega)}$$

$$= \log \left(\frac{\left(\frac{1}{1+\exp(-x^T \omega)} \right)}{\left(\frac{1}{1+\exp(x^T \omega)} \right)} \right)$$

$$= \log \left(\frac{1+\exp(x^T \omega)}{1+\exp(-x^T \omega)} \right)$$

$$= \log \left(\frac{1+\exp(x^T \omega)}{(1+\exp(x^T \omega))/\exp(x^T \omega)} \right)$$

$$= \underline{\underline{x^T \omega}}$$

we want margin > 0 for correct classification

\therefore

if $y = 1$

then $y \cdot (x^T w) > 0$
 $\Rightarrow x^T w > 0$

if $y = -1$
 then

$y \cdot (x^T w) > 0$
 $\Rightarrow x^T w < 0$

from these
2

cond-s we
can see that

$$x^T w = 0$$

This eqⁿ is scale
invariant

$$x^T (cw) = c(x^T w) = 0$$

\updownarrow

$$x^T w = 0$$

$\forall c \in \mathbb{R}$

\rightarrow is the decision
boundary of log.
reg.

Ans 19:

Ans 19)

$$\log L(\hat{c}\hat{w}) = - \sum_{i=1}^n \left[\left(\frac{1+y^{(i)}}{2} \right) \log (1 + \exp(-x^{(i)\top}(\hat{c}\hat{w}))) + \left(\frac{1-y^{(i)}}{2} \right) \log (1 + \exp(x^{(i)\top}(\hat{c}\hat{w}))) \right]$$

$$= - \sum_{i=1}^n \log (1 + \exp(y^{(i)} x^{(i)\top}(\hat{c}\hat{w})))$$

$$\begin{aligned} \frac{\partial}{\partial c} \log L(\hat{c}\hat{w}) &= - \sum_{i=1}^n \frac{\partial}{\partial c} \left[\log (1 + \exp(y^{(i)} x^{(i)\top}(\hat{c}\hat{w}))) \right] \\ &= - \sum_{i=1}^n \frac{-y^{(i)} x^{(i)\top} \hat{w}}{1 + \exp(y^{(i)} x^{(i)\top}(\hat{c}\hat{w}))} \end{aligned}$$

$$\boxed{\frac{\partial}{\partial c} \log(L(\hat{c}\hat{w})) = \sum_{i=1}^n \frac{y^{(i)} x^{(i)\top} \hat{w}}{1 + \exp(y^{(i)} x^{(i)\top}(\hat{c}\hat{w}))}}$$

for linearly separable data-

$$y^{(i)} x^{(i)T} \hat{w} > 0$$

$$\therefore \frac{\partial}{\partial c} \log(L(c\hat{w})) > 0$$

\therefore The log likelihood (and therefore the likelihood) always increases as c increases

Therefore MLE is not well defined in this case as \hat{w} can be scaled to make the likelihood arbitrarily large

Ans 20:

Regularized Logistic Regression

Ans 20

$$J_{\text{logistic}}(w) = \underbrace{\frac{1}{n} \sum_{i=1}^n \log(1 + \exp(-y^{(i)} w^T x^{(i)}))}_{J_1} + \underbrace{\lambda \|w\|_2^2}_{J_2}$$

$$J_{\text{logistic}} : \mathbb{R}^n \rightarrow \mathbb{R}.$$

dom $J_{\text{logistic}} \equiv \mathbb{R}^n$ and is a convex set.
 checking convexity of $J_1 \Rightarrow$

$$\begin{aligned} J_1(w) &= \log(1 + \exp(-y w^T x)) \\ &= \log(1 + \exp(-y x^T w)) \end{aligned}$$

$$\text{let } -y x^T w = t, \quad (t \in \mathbb{R})$$

$$\Rightarrow J(t) = \log(1 + \exp(t))$$

$$J(t) = \log(1 + \exp(t))$$

→ $\exp(t)$ is a convex function
 → 1 is a convex function.

$$J(t) = w_1(1) + w_2(\exp(t))$$

$$[w_1=1, w_2=1]$$

(from 3.2.1 || Rosenberg's notes on optimization)

∴ $J(t)$ is a convex function.

⇒ $J_1(w) = \log(1 + \exp(-y w^T x))$ is convex.

$J_2(w) = \lambda \|w\|_2^2$ is convex $[\lambda \geq 0]$

(Every norm on \mathbb{R}^n is convex)

($J_2(w) = w_1 \|w\|_2^2$ $[w_1 = \lambda, w_1 \geq 0]$)

$$\begin{aligned}
 J(\omega) &= \frac{1}{n} \left(\sum_{i=1}^n (1) (J_1^{\hat{u}_i}(\omega)) + (1) (J_2^{\hat{u}_i}(\omega)) \right) \\
 &= \left(\frac{1}{n} \right) J_1^1(\omega) + \left(\frac{1}{n} \right) J_1^2(\omega) \dots + \left(\frac{1}{n} \right) J_1^n(\omega) \\
 &\quad + \left(\frac{1}{n} \right) J_2^1(\omega) + \left(\frac{1}{n} \right) J_2^2(\omega) \dots + \left(\frac{1}{n} \right) J_2^n(\omega)
 \end{aligned}$$

[where $J_1^{\hat{u}_i}(\omega)$ and $J_2^{\hat{u}_i}(\omega)$ have
 $x = x^{\hat{u}_i}$ and $y = y^{\hat{u}_i}$]

$1/n \geq 0$ as $n \geq 0$

Therefore $J(\omega)$ is a convex function.

(from 3.2.1 || Rosenberg's notes on optimization)

Ans 21:

```
def f_objective(theta, X, Y, l2_param=1):  
    '''  
    Args:  
        theta: 1D numpy array of size num_features  
        X: 2D numpy array of size (num_instances, num_features)  
        y: 1D numpy array of size num_instances  
        l2_param: regularization parameter  
  
    Returns:  
        objective: scalar value of objective function  
    '''  
  
    n = X.shape[0]  
  
    z = X @ theta  
  
    logistic_loss = (1/n) * np.sum(np.logaddexp(0, -Y * z))  
  
    regularization_loss = l2_param * (theta.T @ theta)  
  
    return logistic_loss + regularization_loss
```

please turn over

Ans 22:

```
def fit_logistic_reg(X, y, objective_function, l2_param=1):
    '''
    Args:
        X: 2D numpy array of size (num_instances, num_features)
        y: 1D numpy array of size num_instances
        objective_function: function returning the value of the objective
        l2_param: regularization parameter

    Returns:
        optimal_theta: 1D numpy array of size num_features
    '''
    theta = np.zeros(X.shape[1])
    f = lambda i : objective_function(i, X, y, l2_param)

    optimizer_obj = minimize(f, theta)
    return optimizer_obj.x
```

```
X_train, y_train, X_val, y_val = load_data()

X_train, X_val = std_scaler(X_train, X_val)

X_train = np.hstack((X_train, np.ones((X_train.shape[0], 1)))) # Add bias
term
X_val = np.hstack((X_val, np.ones((X_val.shape[0], 1))))

y_train = np.where(y_train == 0, -1, y_train)
y_val = np.where(y_val == 0, -1, y_val)

print(f'X_train.shape: {X_train.shape}')
print(f'y_train.shape: {y_train.shape}')
print(f'X_val.shape: {X_val.shape}')
print(f'y_val.shape: {y_val.shape}')

theta = fit_logistic_reg(X_train, y_train, f_objective, 0.01)

print(f'theta.shape: {theta.shape}')
print(f'classification error - training set: {classification_error(X_train,
y_train, theta)}')
print(f'classification error - validation set: {classification_error(X_val,
y_val, theta)}')
```




```
X_train.shape: (1600, 21)
y_train.shape: (1600,)
X_val.shape: (400, 21)
y_val.shape: (400,)
theta.shape: (21,)
classification error - training set: 0.255
classification error - validation set: 0.26
```

please turn over

Ans 23:

```
def plot_lambda_reg_vs_log_likelihood(X_train, y_train, X_test, y_test, \
    lambda_reg = [1e-7, 1e-6, 1e-5, 1e-3, 1e-2, 1e-1, 1e0, 1e1, 1e2]):
    log_likelihoods = []
    for l in lambda_reg:
        theta = fit_logistic_reg(X_train, y_train, f_objective, l)
        log_likelihoods.append(log_likelihood(theta, X_test, y_test))
    plt.figure(figsize=(10, 10))
    plt.plot([np.log10(i) for i in lambda_reg], log_likelihoods, 'bo-', \
        label = 'log-lambda values vs log likelihoods')
    plt.xlabel('log10-lambda')
    plt.ylabel('NLL')
    plt.show()
```

```
plot_lambda_reg_vs_log_likelihood(X_train, y_train, X_val, y_val)
```




please turn over

Ans 24:

```
def plot_predicted_prob_calibration(theta, X, y, bins = 5):
    y_pred_prob = sigmoid(X @ theta)
    coupled_mat = np.column_stack((y, y_pred_prob))
    coupled_mat = coupled_mat[coupled_mat[:, 1].argsort()[::-1]]
    range_start = 0
    range_end = 0
    probability_ranges = []
    percentage_positive_samples = []
    for i in np.linspace(0, coupled_mat.shape[0], bins, dtype=int):
        if i == range_end:
            continue
        range_start = range_end
        range_end = i
        mat = coupled_mat[range_start:range_end, :]
        probability_ranges.append(f"{np.min(mat[:, 1]):.2f} to
{np.max(mat[:, 1]):.2f}")
        percentage_positive_samples.append((np.sum(mat[:, 0] == 1)) /
mat.shape[0])

    plt.figure(figsize=(10, 10))
    plt.plot(probability_ranges, percentage_positive_samples, 'ro-',
label='score ranges vs classification erros')
    plt.xlabel("predicted probability ranges")
    plt.ylabel("positive examples ratio")
    plt.xticks(fontsize=7)
    plt.yticks(fontsize=7)
    plt.show()
```

```
plot_predicted_prob_calibration(X= X_val, y= y_val, theta= theta, bins= 5)
```



Observations

- The % of positive examples in each prediction probability range changes almost linearly with the mean of the predicted probability ranges having value similar to the % of positive labelled samples in that class
- $\text{sigmoid}((W.T)(X)) \sim (\text{number of positive samples in class} / \text{total number of samples in class})$