

Homework 3: Bayesian ML and Multiclass

Bayesian logistic regression with Gaussian Prior

Bayesian logistic Regression with Gaussian prior.

→ binary classification setting

⇒ input space: $x \in \mathbb{R}^d$

outcome space: $y \in \{-1, 1\}$

dataset: $D = \{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(N)}, y^{(N)})\}$

$$p(y=1|x;w) = \frac{1}{1 + \exp(-x^T w)}$$

for logistic reg. in Bayesian setting.
introduce prior $p(w)$ on $w \in \mathbb{R}^d$

Ans. 1

① posterior density:

$$p(w|D) \propto p(D|w) \cdot p(w)$$

$$\Rightarrow p(w|D) = \kappa \cdot \underbrace{p(D|w)}_{\text{likelihood}} \cdot \underbrace{p(w)}_{\text{prior}}$$

$$\Rightarrow p(w|D) = \kappa \left(\exp(-NLL_D(w)) \right) \cdot p(w) \quad \left(\text{where } NLL \text{ is the negative log likelihood} \right)$$

please turn over

Ans 2.

(2) let $w \sim N(0, \Sigma)$

MAP estimate for w after observing data D

$$\begin{aligned} \arg \max_w p(w|D) &= \arg \max_w p(D|w) \cdot p(w) \\ &= \arg \max_w \left[\log(p(D|w)) + \log(p(w)) \right] \quad \begin{array}{l} \text{[taking log as it} \\ \text{is strictly inc. f'n} \\ \text{\& makes compute easy]} \end{array} \end{aligned}$$

log likelihood function
for log. reg.

$$\arg \max_w \left[\sum_{i=1}^n -\frac{1}{2} (y^{(i)} \log(1 + \exp(-x^{(i)T} w)) + (1 - y^{(i)}) \log(1 + \exp(x^{(i)T} w))) + \log(p(w)) \right]$$

$$= \arg \max_w \left(-NLL_D(w) + \log|2\pi\Sigma|^{-\frac{1}{2}} - \frac{1}{2} w^T \Sigma^{-1} w \right)$$

↳ can be removed as it is
not dep. on w

$$= \arg \min_w \left(NLL_D(w) + \frac{1}{2} w^T \Sigma^{-1} w \right) \equiv \arg \min_w \left(\frac{1}{n} NLL_D(w) + \frac{1}{2n} w^T \Sigma^{-1} w \right)$$

①

L_2 regularized logistic regression

$$= \arg \min_w \left(\frac{1}{n} \sum_{i=1}^n \log(1 + \exp(-y^{(i)} x^{(i)T} w)) + \lambda \|w\|_2^2 \right)$$

↳ can be rewritten as

$$\frac{1}{2n} \left[\sum_{i=1}^n y^{(i)} \log(1 + \exp(-x^{(i)T} w)) + (1 - y^{(i)}) \log(1 + \exp(x^{(i)T} w)) \right]$$

$$= \frac{1}{n} NLL_D(w)$$

$$\Rightarrow \arg \min_w \left(\frac{1}{n} NLL_D(w) + \lambda \|w\|_2^2 \right)$$

②

Equating eq's (2) and (1) we see that

$$\frac{1}{2n} (w^T \Sigma^{-1} w) = \lambda \|w\|_2^2$$

$$\Rightarrow \Sigma = \frac{1}{2n\lambda} I$$

please turn over

Ans 3.

③ for ERM to be the mode of the posterior dist.
such that $\Sigma = I$.

$$\Rightarrow I = \frac{1}{2n\lambda} I$$

$$\Rightarrow \lambda = \frac{1}{2n}$$

please turn over

Coin Flipping with partial observability

Coin flipping with partial observability.

$$p(z = H \mid \theta_1) = \theta_1 \quad // \text{biased coin}.$$

someone reports z to us as x
there is a chance x is incorrect if $z = H$.
this is denoted by

$$\text{and } p(x = H \mid z = H, \theta_2) = \theta_2$$

$$p(x = T \mid z = T) = 1.$$

please turn over

Ans 4.

④ To prove $p(x=H | \theta_1, \theta_2) = \theta_1 \theta_2$. ①

Proof:
$$p(x=H | \theta_1, \theta_2) = \left[\underbrace{p(x=H | z=H, \theta_1, \theta_2) \cdot p(z=H | \theta_1)}_{\text{②}} + \underbrace{p(x=H | z=T, \theta_1, \theta_2) \cdot p(z=T | \theta_1)}_{\text{②}} \right]$$

Evaluating ①

$$\Rightarrow p(x=H | z=H, \theta_1, \theta_2) \cdot p(z=H | \theta_1)$$

↳ since it is known that $z=H$ for this case

$$\therefore p(x=H | z=H, \theta_1, \theta_2) = p(x=H | z=H, \theta_2) = \theta_2$$

and $p(z=H | \theta_1) = \theta_1$

$$\therefore \text{①} = \theta_1 \theta_2$$

Evaluating ②

$$p(x=H | z=T, \theta_1, \theta_2) = p(z=T | \theta_1)$$

↳ since x always correctly reports z if $z=T$

$$\therefore p(x=H | z=T, \theta_1, \theta_2) = 0$$

and $p(z=T | \theta_1) = (1 - \theta_1)$

$$\left. \begin{array}{l} p(x=H | z=T, \theta_1, \theta_2) = 0 \\ p(z=T | \theta_1) = (1 - \theta_1) \end{array} \right\} \therefore \text{②} = 0$$

$$\therefore \boxed{p(x=H | \theta_1, \theta_2) = \text{①} + \text{②} = \theta_1 \theta_2.} \quad // \text{ Hence proved}$$

please turn over

Ans 5.

⑤ assuming all events are independent, we get -

$$\mathcal{L}_{D_N}(\theta_1, \theta_2) = \prod_{i=1}^{N_H} p(x_i | \theta_1, \theta_2)$$

up. incorrectly reported head + all tails correctly reported.

$$\mathcal{L}_{D_N}(\theta_1, \theta_2) = (\theta_1, \theta_2)^{n_H} \overline{(1 - \theta_1, \theta_2)^{n_T}}$$

$\overline{\text{correctly reported H}}$

please turn over

Ans 6.

⑥ Estimating θ_1 and θ_2 through MLE is not usable as say for a given distribution, we get

$$\chi_{D_n}(\theta_1, \theta_2) = (\theta_1 \theta_2)^{n_h} (1 - \theta_1 \theta_2)^{n_t}$$

if we try MLE, we get

$$\arg \max_{\theta_1, \theta_2} \log \chi_{D_n}(\theta_1, \theta_2)$$

$$= \arg \max_{\theta_1, \theta_2} n_h \log \theta_1 \theta_2 + n_t \log (1 - \theta_1 \theta_2)$$

$$\frac{\partial \log \chi(\theta_1, \theta_2)}{\partial \theta_1} = 0 = \frac{n_h}{\theta_1 \theta_2} \cdot \theta_2 + \frac{n_t}{1 - \theta_1 \theta_2} \cdot (-\theta_2) \quad \left. \begin{array}{l} \theta_1, \theta_2 \in (0, 1] \\ \text{[for this to be differentiable]} \end{array} \right\}$$

$$\Rightarrow \theta_2 (n_h - n_h \theta_1 \theta_2) - n_t \theta_1 \theta_2^2 = 0$$

$$\Rightarrow n_h \theta_2 - n_h \theta_1 \theta_2^2 - n_t \theta_1 \theta_2^2 = 0$$

$$\Rightarrow (\theta_2) (n_h - \theta_1 \theta_2 (n_h + n_t)) = 0$$

$$\theta_1 \theta_2 = \frac{n_h}{n_h + n_t} \quad \left. \begin{array}{l} \text{ } \\ \text{ } \end{array} \right\}$$

$$\frac{\partial \log \chi(\theta_1, \theta_2)}{\partial \theta_2} = 0 = \frac{n_h}{\theta_1 \theta_2} \theta_1 + \frac{n_t}{1 - \theta_1 \theta_2} \cdot (-\theta_1)$$

$$\Rightarrow (\theta_1) (n_h - (n_h + n_t) \theta_1 \theta_2) = 0$$

$$\Rightarrow \theta_1 \theta_2 = \frac{n_h}{n_h + n_t} \quad \left. \begin{array}{l} \text{ } \\ \text{ } \end{array} \right\}$$

identical
∴ we have
2 variables
1 eqⁿ
non trivial solⁿ only
possible.

Hence since we get non trivial solⁿs for θ_1 and θ_2 we cannot estimate θ_1 and θ_2 using MLE.

please turn over

Ans 7.

⑦ Assuming all events are independent, we can write the likelihood as

$$L(\theta_1, \theta_2) = p(D_H, D_C | \theta_1, \theta_2) = p(D_H | \theta_1, \theta_2) \cdot p(D_C | \theta_1, \theta_2)$$

$$= (\theta_1 \theta_2)^{n_H} (1 - \theta_1 \theta_2)^{n_T} (\theta_1)^{c_H} (1 - \theta_1)^{c_T}$$

$$\text{arg max}_{\theta_1, \theta_2} (\log L(\theta_1, \theta_2)) \Rightarrow$$

$$l(\theta_1, \theta_2) = n_H \log(\theta_1 \theta_2) + n_T \log(1 - \theta_1 \theta_2) + c_H \log \theta_1 + c_T \log(1 - \theta_1)$$

Differentiating wrt θ_1 we get:-

$$\frac{\partial l}{\partial \theta_1} = \frac{n_H}{\theta_1 \theta_2} + \frac{n_T (-\theta_2)}{1 - \theta_1 \theta_2} + \frac{c_H}{\theta_1} + \frac{c_T (-1)}{1 - \theta_1} = 0$$

$$= \frac{\theta_2}{(\theta_1 \theta_2)(1 - \theta_1 \theta_2)} (n_H - \theta_1 \theta_2 (n_H + n_T)) + \frac{1}{\theta_1(1 - \theta_1)} (c_H - \theta_1 (c_H + c_T)) = 0$$

$$\Rightarrow 0 = \frac{(n_H - (\theta_1 \theta_2)(n_H + n_T))}{\theta_1(1 - \theta_1 \theta_2)} + \frac{(c_H - \theta_1(c_H + c_T))}{\theta_1(1 - \theta_1)} \quad \text{--- (1)}$$

Differentiating wrt θ_2 :-

$$\frac{\partial l}{\partial \theta_2} = 0 = \frac{(n_H - (\theta_1 \theta_2)(n_H + n_T))}{\theta_2(1 - \theta_1 \theta_2)} \quad \text{--- (2)}$$

solving these eq's 1 & 2 we get

$$\theta_1 = \frac{c_H}{c_H + c_T}$$

and

$$\theta_2 = \frac{n_H}{n_H + n_T} \cdot \frac{c_H + c_T}{c_H}$$

non trivial solution.
∴ we can estimate θ_1, θ_2 using MLE.

please turn over

Ans 8.

8

$$\theta_1 \sim \text{Beta}(h, t)$$

$$\Rightarrow p(\theta_1) = \theta_1^{h-1} (1-\theta_1)^{t-1}$$

prior is provided to mitigate the overfitting in D_c

\Rightarrow posterior of θ_1 based on clean results D_c

$$p(\theta_1 | D_c) \propto p(D_c | \theta_1) \cdot p(\theta_1)$$

$$\Rightarrow p(\theta_1 | D_c) = K \cdot \theta_1^{c_h} (1-\theta_1)^{c_t} \cdot \theta_1^{h-1} (1-\theta_1)^{t-1}$$

$$= K \theta_1^{c_h+h-1} (1-\theta_1)^{c_t+t-1}$$

$$p(\theta_1 | D_c) = \text{Beta}(c_h+h, c_t+t)$$

$$\therefore \arg \max_{\theta_1} p(\theta_1 | D_c) = \frac{c_h+h-1}{c_h+h+c_t+t-2} = \theta_{1 \text{ MAP}}$$

$$\therefore \theta_2 = \frac{n_h}{n_h+n_t} \cdot \frac{c_h+h+c_t+t-2}{c_h+h-1}$$

↳ formula in slides.

↳ from eqⁿ ② of problem 7]

please turn over

ℓ_2 -regularized empirical risk function for multiclass hinge loss

Ans 9.

⑨ To prove : $J(w)$ is convex function of w .

Proof: →

$$\text{Let } f_i(w) = \Delta C(y_i, y) + (\langle w, \psi(x_i, y) - \psi(x_i, y_i) \rangle)$$

$$\left. \begin{aligned} \langle w, \psi(x_i, y) \rangle &= \langle w, \psi(x_i, y) \rangle + 0 \\ \Delta C(y_i, y) &\in \mathbb{R} \end{aligned} \right\} \text{affine function}$$

$$\langle w, \psi(x_i, y_i) \rangle \in \mathbb{R}.$$

∴ $f_i(w)$ is an affine function of y on w and thus convex.

$$\Rightarrow \max_{y \in Y} f_y(w) \equiv \max \{ f_1(w), f_2(w), \dots, f_K(w) \}$$

is also convex with domain
 $\text{dom } f_1 \cap \text{dom } f_2 \cap \dots \cap \text{dom } f_K$

→ (Rosenberg's
 notes
 3.2.4)

$$\Rightarrow \frac{1}{n} \sum_{i=1}^n \max_{y \in \mathcal{Y}} f_y(w) \text{ is convex } \rightarrow (\text{Rosenberg's notes 3.2.1})$$

Now $\lambda \|w\|_2^2$ is convex as $\lambda \in \mathbb{R}$ and $\lambda > 0$ and $\|w\|_2^2$ is convex
 \rightarrow (Rosenberg's notes 3.1.3)

∴ $\lambda \|w\|_2^2 + \frac{1}{n} \sum_{i=1}^n \max_{y \in \mathcal{Y}} f_y(w)$ is convex [sum of 2 convex f's is convex]

$$\Rightarrow \lambda \|w\|_2^2 + \frac{1}{n} \sum_{i=1}^n \max_{y \in \mathcal{Y}} \left[\Delta(y_i, y) + \langle w, \psi(x_i, y) - \psi(x_i, y_i) \rangle \right]$$

is a convex function.

hence Proved.

please turn over

Ans 10.

10

$$J(w) = \lambda \|w\|_2^2 + \frac{1}{n} \sum_{i=1}^n \max_{y \in Y} [\Delta(y_i, y) + \langle w, \psi(x_i, y) - \psi(x_i, y_i) \rangle]$$

Let $R(w) = \lambda \|w\|_2^2$ has subgradient ∇

$$g_R(w) = 2\lambda w \in \partial R(w)$$

$$\text{Let } H(w) = \frac{1}{n} \sum_{i=1}^n \max_{y \in Y} [\Delta(y_i, y) + \langle w, \psi(x_i, y) - \psi(x_i, y_i) \rangle]$$

for each i
$$\hat{y}_i = \arg \max_{y \in Y} [\Delta(y_i, y) + \langle w, \psi(x_i, y) - \psi(x_i, y_i) \rangle]$$

$\hookrightarrow \hat{y}_i$ is the y that maximizes above exp.

from Assignment 2 we know that

$$\text{given } f(x) = \max_k f_k(x) \quad (k = \{1, 2, \dots, n\})$$

if k is any index for which

$$f_k(x) = f(x)$$

$$\text{then if } g(x) \in \partial f_k(x)$$

$$\Rightarrow g(x) \in \partial f(x)$$

$$\text{Let } f_y(w) = \Delta(y_i, y) + \langle w, \psi(x_i, y) - \psi(x_i, y_i) \rangle$$

\therefore subgradient of $f_y(w)$ at w can be given by subgradient of $f_{\hat{y}_i}(w)$

$$\text{where } \hat{y}_i = \arg \max_{y \in Y} f_y(w)$$

one subgradient of $f_{\hat{y}_i}(w)$ wrt w is \Rightarrow

$$g_{f_{\hat{y}_i}}(w) = \psi(x_i, \hat{y}_i) - \psi(x_i, y_i) \in \partial f(w)$$

\therefore combining all components of $H(w)$ we get

$$g_H(w) = \frac{1}{n} \sum_{i=1}^n (\psi(x_i, \hat{y}_i) - \psi(x_i, y_i))$$

we can write

$$g_J(w) = g_R(w) + g_H(w)$$

$$\Rightarrow g_J(w) = 2\lambda w + \frac{1}{n} \sum_{i=1}^n (\psi(x_i, \hat{y}_i) - \psi(x_i, y_i^*))$$

$\left\{ \text{where } \hat{y}_i = \arg \max_{y \in Y} [\lambda C(y_i, y) + \langle w, \phi(x_i, y) - \phi(x_i, y_i^*) \rangle] \right\}$

$$g_J(w) \in \partial C_J(w)$$

please turn over

Ans 11.

⑪ stochastic gradient descent will be based on a single training eq. (x_i, y_i)

$$\Rightarrow \nabla \left(J(w, x_i^o, y_i^o) \right) = \left(2\lambda w + \left(\psi(x_i, \hat{y}_i) - \psi(x_i, y_i^o) \right) \right)$$

please turn over

Ans 12.

12 minibatch gradient will be based on avg of samples in the mini-batch.

$$\partial \left(J(w, (x_i, y_i), (x_{i+1}, y_{i+1}) \dots (x_{i+m-1}, y_{i+m-1})) \right)$$

||

$$\left(2\lambda w + \frac{1}{m} \sum_{t=i}^{i+m-1} \left(\psi(x_t, \hat{y}_t) - \psi(x_t, y_t) \right) \right)$$

please turn over

Hinge Loss is a specialized case of Generalized Hinge Loss

Let $\mathcal{Y} = \{-1, 1\}$. Let $\Delta(y, \hat{y}) = \mathbb{1}_{y \neq \hat{y}}$. If $g(x)$ is the score function in our binary classification setting, then define our compatibility function as

$$\begin{aligned} h(x, 1) &= g(x)/2 \\ h(x, -1) &= -g(x)/2. \end{aligned}$$

Show that for this choice of h , the multiclass hinge loss reduces to hinge loss:

$$\ell(h, (x, y)) = \max_{y' \in \mathcal{Y}} [\Delta(y, y') + h(x, y') - h(x, y)] = \max\{0, 1 - yg(x)\}$$

In this problem we will work on a simple three-class classification example. The data is generated and plotted for you in the skeleton code.

Hinge loss is a Specialized Case of Generalized Hinge loss

$$\Delta(y, \hat{y}) = 1_{y \neq \hat{y}} = \begin{cases} 0, & y = \hat{y} \\ 1, & y \neq \hat{y} \end{cases}$$

Compatibility function

$$h(x, 1) = g(x)/2, \quad h(x, -1) = -g(x)/2.$$

generalized multiclass hinge loss

$$L(h, (x, y)) = \max_{y' \in Y} [\Delta(y, \hat{y}) + h(x, y') - h(x, y)]$$

Case 1: correct classification ($y' = y$)

$$\Delta(y, y') = 0$$

$$h(x, y') = h(x, y)$$

$$\Rightarrow \Delta(y, y') + h(x, y') - h(x, y) = 0$$

Case 2: incorrect classification. ($\hat{y} \neq y$)

$$\Rightarrow \Delta(y, \hat{y}) = 1$$

Case 2a: $y = 1, \hat{y} = -1$

$$\Rightarrow \Delta(y, \hat{y}) + h(x, y') - h(x, y)$$

$$= 1 + \frac{-g(x)}{2} - \frac{g(x)}{2}$$

$$= 1 - g(x)$$

Case 2b: $y = -1, \hat{y} = 1$

$$\Rightarrow \Delta(y, \hat{y}) + h(x, y') - h(x, y)$$

$$= 1 + g(x).$$

We can combine cases 2a and 2b to get a gen. expr. for case 2 as follows:

$$\text{Case 2:} \\ = 1 - y \cdot g(x)$$

∴

$$\begin{aligned} d(h, (x, y)) &= \max_{y' \in Y} [\Delta(y, y') + h(x, y') - h(x, y)] \\ &= \max(0, 1 - y \cdot g(x)) \end{aligned} \quad \left. \vphantom{\begin{aligned} d(h, (x, y)) &= \max_{y' \in Y} [\Delta(y, y') + h(x, y') - h(x, y)] \\ &= \max(0, 1 - y \cdot g(x)) \end{aligned}} \right\} \text{taking the max. of} \\ &\quad \text{case 1 and case 2}$$

Hence proved.

One-vs-All (also known as One-vs-Rest)

Ans 13.

- ```
def fit(self, X, y=None):
 #Your code goes here
 for i in range(self.n_classes):
 class_i_labels = np.where(y == i, 1, 0)
 self.estimators[i].fit(X, class_i_labels)
 self.fitted = True
 return self
```
- ```
def decision_function(self, X):
    if not self.fitted:
        raise RuntimeError("You must train classifier \
before predicting data.")

    if not hasattr(self.estimators[0], "decision_function"):
        raise AttributeError(
            "Base estimator doesn't have a\
decision_function attribute.")
    scores = np.empty((X.shape[0], self.n_classes))
    for i in range(self.n_classes):
        class_i_predicted_labels =
self.estimators[i].decision_function(X)
        scores[:, i] = class_i_predicted_labels
    return scores
```
- ```
def predict(self, X):
 """
 Predict the class with the highest score.
 @param X: array-like, shape = [n_samples,n_features] input data
 @returns array-like, shape = [n_samples,] the predicted classes
 for each input
 """
 scores = self.decision_function(X)
 return np.argmax(scores, axis = 1)
```

please turn over

Ans 14.

```
#Here we test the OneVsAllClassifier
from sklearn import svm
svm_estimator = svm.LinearSVC(loss='hinge', fit_intercept=False,
C=200)
clf_onevsall = OneVsAllClassifier(svm_estimator, n_classes=3)
clf_onevsall.fit(X,y)

for i in range(3) :
 print("Coeffs %d"%i)
 print(clf_onevsall.estimators[i].coef_) #Will fail if you
haven't implemented fit yet


create a mesh to plot in
h = .02 # step size in the mesh
x_min, x_max = min(X[:,0])-3,max(X[:,0])+3
y_min, y_max = min(X[:,1])-3,max(X[:,1])+3
xx, yy = np.meshgrid(np.arange(x_min, x_max, h),
 np.arange(y_min, y_max, h))
mesh_input = np.c_[xx.ravel(), yy.ravel()]

Z = clf_onevsall.predict(mesh_input)
Z = Z.reshape(xx.shape)
plt.contourf(xx, yy, Z, cmap=plt.cm.coolwarm, alpha=0.8)
Plot also the training points
plt.scatter(X[:, 0], X[:, 1], c=y, cmap=plt.cm.coolwarm)

from sklearn import metrics
metrics.confusion_matrix(y, clf_onevsall.predict(X))
```

```
Coeffs 0
[[-1.05853334 -0.90294603]]
Coeffs 1
[[0.42121645 0.27171776]]
Coeffs 2
[[0.89164752 -0.82601734]]
/home/arjun-prasad/anaconda3/envs/deep_learning/lib/python3.12/site-packages/sklearn/svm/_base
warnings.warn(
```

```
array([[100, 0, 0],
 [0, 100, 0],
 [0, 11, 89]])
```



please turn over



## Multiclass SVM

Ans 15.

- ```
def featureMap(X,y,num_classes) :  
    #The following line handles X being a 1d-array or a 2d-array  
    num_samples, num_inFeatures = (1,X.shape[0]) if \  
        len(X.shape) == 1 else (X.shape[0],X.shape[1])  
    #your code goes here, and replaces following return  
    if X.ndim == 1:  
        X.reshape(1, -1)  
  
    featureMap = np.zeros((num_samples, num_classes*num_inFeatures))  
    loc = num_inFeatures*y  
    for i in range(num_samples):  
        featureMap[i, loc : loc + num_inFeatures] = X[i]  
  
    return featureMap
```

please turn over

Ans 16.

- ```
def sgd(X, y, num_outFeatures, subgd, eta = 0.1, T = 10000):
 num_samples = X.shape[0]
 #your code goes here and replaces following return statement
 w = np.zeros(num_outFeatures)
 for iter in range(T):
 idx = np.random.randint(0, num_samples)
 grad = subgd(X[idx], y[idx], w)
 w = w - (eta*grad)
 return w
```

please turn over

Ans 17.

- ```
def subgradient(self,x,y,w):
    """
    Computes the subgradient at a given data point x,y
    @param x: sample input
    @param y: sample class
    @param w: parameter vector
    @return returns subgradient vector at given x,y,w
    """

    #Your code goes here and replaces the following return statement
    subgrad_R = 2*self.lam*w
    y_hat = -1

    f_y_w_max = -np.inf
    for yi in range(self.num_classes):
        if yi == y:
            continue # no point in evaluating this
        f_yi_w = self.Delta(yi, y) + (w @ (self.Psi(x, yi).flatten()
- \
            self.Psi(x, y).flatten()))
        if f_yi_w > f_y_w_max:
            f_y_w_max = f_yi_w
            y_hat = yi
    return subgrad_R + (self.Psi(x, y_hat).flatten() - self.Psi(x,
y).flatten())
```

- ```
def decision_function(self, X):
 if not self.fitted:
 raise RuntimeError("You must train classifier before
predicting data.")

 #Your code goes here and replaces following return statement
 num_samples = X.shape[0]
 class_scores = np.zeros((num_samples, self.num_classes))
 for yi in range(self.num_classes):
 class_scores[:, yi] = self.Psi(X, yi) @ self.coef_
 return class_scores
```

- ```
def predict(self, X):
    #Your code goes here and replaces following return statement
    class_scores = self.decision_function(X)
    return np.argmax(class_scores, axis = 1)
```

please turn over

Ans 18.

- ```

from skeleton_code import zeroOne, featureMap, sgd, MulticlassSVM
#the following code tests the MulticlassSVM and sgd
#will fail if MulticlassSVM is not implemented yet
est = MulticlassSVM(6, lam=1)
est.fit(X,y)
print("w:")
print(est.coef_)
Z = est.predict(mesh_input)
Z = Z.reshape(xx.shape)
plt.contourf(xx, yy, Z, cmap=plt.cm.coolwarm, alpha=0.8)
Plot also the training points
plt.scatter(X[:, 0], X[:, 1], c=y, cmap=plt.cm.coolwarm)

from sklearn import metrics
metrics.confusion_matrix(y, est.predict(X))

```

```

w:
[-0.41923524 -0.41923524 -0.06249277 -0.06249277 0.48172801 0.48172801]

array([[100, 0, 0],
 [0, 100, 0],
 [0, 0, 100]])

```

