

Homework 4: Decision Trees, Boosting, and Neural Networks

Decision Tree Implementation

Ans. 1

```
def compute_entropy(label_array):  
    '''  
    Calulate the entropy of given label list  
  
    :param label_array: a numpy array of binary labels shape = (n, 1)  
    :return entropy: entropy value  
    '''  
    # count the different labels and their counts in the label array  
    label_val, label_count = np.unique(label_array, return_counts= True)  
    tot_count = np.sum(label_count)  
    label_prob = label_count / tot_count  
    epsilon = 1e-12  
    entropy = -(label_prob * np.log(label_prob + epsilon)).sum()  
    return entropy  
  
def compute_gini(label_array):  
    '''  
    Calulate the gini index of label list  
  
    :param label_array: a numpy array of labels shape = (n, 1)  
    :return gini: gini index value  
    '''  
    label_val, label_count = np.unique(label_array, return_counts= True)  
    tot_count = np.sum(label_count)  
    label_prob = label_count / tot_count  
    gini = (np.sum(label_prob * (1 - label_prob)))  
    return gini
```

please turn over

Ans 2

```

class Decision_Tree(BaseEstimator):

    def __init__(self, split_loss_function, leaf_value_estimator,
                  depth=0, min_sample=5, max_depth=10):
        """
        Initialize the decision tree classifier

        :param split_loss_function: method with args (X, y) returning loss
        :param leaf_value_estimator: method for estimating leaf value from
array of ys
        :param depth: depth indicator, default value is 0, representing
root node
        :param min_sample: an internal node can be splitted only if it
contains points more than min_smaple
        :param max_depth: restriction of tree depth.
        """
        self.split_loss_function = split_loss_function
        self.leaf_value_estimator = leaf_value_estimator
        self.depth = depth
        self.min_sample = min_sample
        self.max_depth = max_depth
        self.is_leaf = False
        """
        ----EDIT: adding additional parameters to the decision tree class--
        """
        self.left = None
        self.right = None
        self.split_id = -1
        self.split_value = None
        self.value = None

    def fit(self, x, y):
        """
        This should fit the tree classifier by setting the values
self.is_leaf,
        self.split_id (the index of the feature we want ot split on, if
we're splitting),
        self.split_value (the corresponding value of that feature where the
split is),
        and self.value, which is the prediction value if the tree is a leaf
node. If we are
        splitting the node, we should also init self.left and self.right to
be Decision_Tree
        objects corresponding to the left and right subtrees. These
subtrees should be fit on
        the data that fall to the left and right, respectively, of
self.split_value.
        This is a recursive tree building procedure.

```

```

        :param X: a numpy array of training data, shape = (n, m)
        :param y: a numpy array of labels, shape = (n, 1)

        :return self
        """
        # first check if the current node is to be a leaf or has to be
        split further
        if((self.depth >= self.max_depth) or (x.shape[0] <
self.min_sample)):
            self.is_leaf = True
            self.value = self.leaf_value_estimator(y)
            return self

        # in case the node needs to be split - we need to create the
        subsets of data that need to be given to each child node
        self.find_best_feature_split(x, y)

        if self.split_id == -1:
            self.is_leaf = True
            self.value = self.leaf_value_estimator(y)
            return self

        left_idx = np.where(x[:, self.split_id] <= self.split_value)
        right_idx = np.where(x[:, self.split_id] > self.split_value)

        if np.sum(left_idx) == 0 or np.sum(right_idx) == 0:
            self.split_id = -1
            self.is_leaf = True
            self.value = self.leaf_value_estimator(y)
            return self

        # create the left and right subtrees
        self.left = Decision_Tree(self.split_loss_function,
self.leaf_value_estimator, self.depth + 1, self.min_sample, self.max_depth)
        self.right = Decision_Tree(self.split_loss_function,
self.leaf_value_estimator, self.depth + 1, self.min_sample, self.max_depth)

        self.left.fit(x[left_idx], y[left_idx])
        self.right.fit(x[right_idx], y[right_idx])

        return self

def find_best_split(self, x_node, y_node, feature_id):
    """
    For feature number feature_id, returns the optimal splitting point
    for data X_node, y_node, and corresponding loss.

    :param x_node: a numpy array of training data, shape = (n_node, m)
    :param y_node: a numpy array of labels, shape = (n_node, 1)
    """
    min_loss = np.inf
    split_value = np.inf

    x_node_ftr = x_node[:, feature_id] # Extract the feature column

```

```

num_samples = len(y_node)

# Sort the input feature and maintain corresponding labels
sort_idx = np.argsort(x_node_ftr)
sorted_x_node_ftr = x_node_ftr[sort_idx]
sorted_y_node = y_node[sort_idx]

# Identify unique values in the sorted feature column
unique_values = np.unique(sorted_x_node_ftr)

for value in unique_values:
    # Split the data based on the unique value
    l_idx_mask = sorted_x_node_ftr <= value
    r_idx_mask = ~l_idx_mask # Complement of the left mask

    y_left = sorted_y_node[l_idx_mask]
    y_right = sorted_y_node[r_idx_mask]

    # Skip if either side is empty
    if len(y_left) == 0 or len(y_right) == 0:
        continue

    # Compute loss
    net_loss = (
        (len(y_left) / num_samples) *
self.split_loss_function(y_left) +
        (len(y_right) / num_samples) *
self.split_loss_function(y_right)
    )

    # Update if current loss is better
    if net_loss < min_loss:
        min_loss = net_loss
        split_value = value

return split_value, min_loss

def find_best_feature_split(self, x_node, y_node):
    """
    Returns the optimal feature to split and best splitting point
    for data X_node, y_node.

    ----EDIT: I have edited the shape of param X to (n_node, m) (prvsly
mentioned as (n_node))----
    :param X: a numpy array of training data, shape = (n_node, m)

    :param y: a numpy array of labels, shape = (n_node, 1)
    """
    min_loss = np.inf
    for i in range(x_node.shape[1]):
        # print(f"find best feature split - feature: {i}")
        curr_feature = i
        curr_feature_min_split_val, curr_feature_min_loss =
self.find_best_split(x_node, y_node, feature_id= curr_feature)

```

```
        if(min_loss > curr_feature_min_loss):
            self.split_id = curr_feature # curr best feature found
            self.split_value = curr_feature_min_split_val # current
value of best feature on which the split loss is minimum
            min_loss = curr_feature_min_loss

    def predict_instance(self, instance):
        """
        Predict label by decision tree

        :param instance: a numpy array with new data, shape (1, m)

        :return whatever is returned by leaf_value_estimator for leaf
containing instance
        """
        if self.is_leaf:
            return self.value
        if instance[self.split_id] <= self.split_value:
            return self.left.predict_instance(instance)
        else:
            return self.right.predict_instance(instance)
```

please turn over

Ans. 3

```

def most_common_label(y):
    """
    Find most common label
    """
    label_cnt = Counter(y.reshape(len(y)))
    label = label_cnt.most_common(1)[0][0]
    return label

class Classification_Tree(BaseEstimator, ClassifierMixin):

    loss_function_dict = {
        'entropy': compute_entropy,
        'gini': compute_gini
    }

    def __init__(self, loss_function='entropy', min_sample=5,
max_depth=10):
        """
        :param loss_function(str): loss function for splitting internal
node
        """

        self.tree = Decision_Tree(self.loss_function_dict[loss_function],
                                most_common_label,
                                0, min_sample, max_depth)

    def fit(self, X, y=None):
        self.tree.fit(X,y)
        return self

    def predict_instance(self, instance):
        value = self.tree.predict_instance(instance)
        return value

# Training classifiers with different depth
clf1 = Classification_Tree(max_depth=1, min_sample=2)
clf1.fit(x_train, y_train_label)

clf2 = Classification_Tree(max_depth=2, min_sample=2)
clf2.fit(x_train, y_train_label)

clf3 = Classification_Tree(max_depth=3, min_sample=2)
clf3.fit(x_train, y_train_label)

clf4 = Classification_Tree(max_depth=4, min_sample=2)
clf4.fit(x_train, y_train_label)

clf5 = Classification_Tree(max_depth=5, min_sample=2)
clf5.fit(x_train, y_train_label)

```

```

clf6 = Classification_Tree(max_depth=6, min_sample=2)
clf6.fit(x_train, y_train_label)

# Plotting decision regions
x_min, x_max = x_train[:, 0].min() - 1, x_train[:, 0].max() + 1
y_min, y_max = x_train[:, 1].min() - 1, x_train[:, 1].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.1),
                     np.arange(y_min, y_max, 0.1))

f, axarr = plt.subplots(2, 3, sharex='col', sharey='row', figsize=(10, 8))

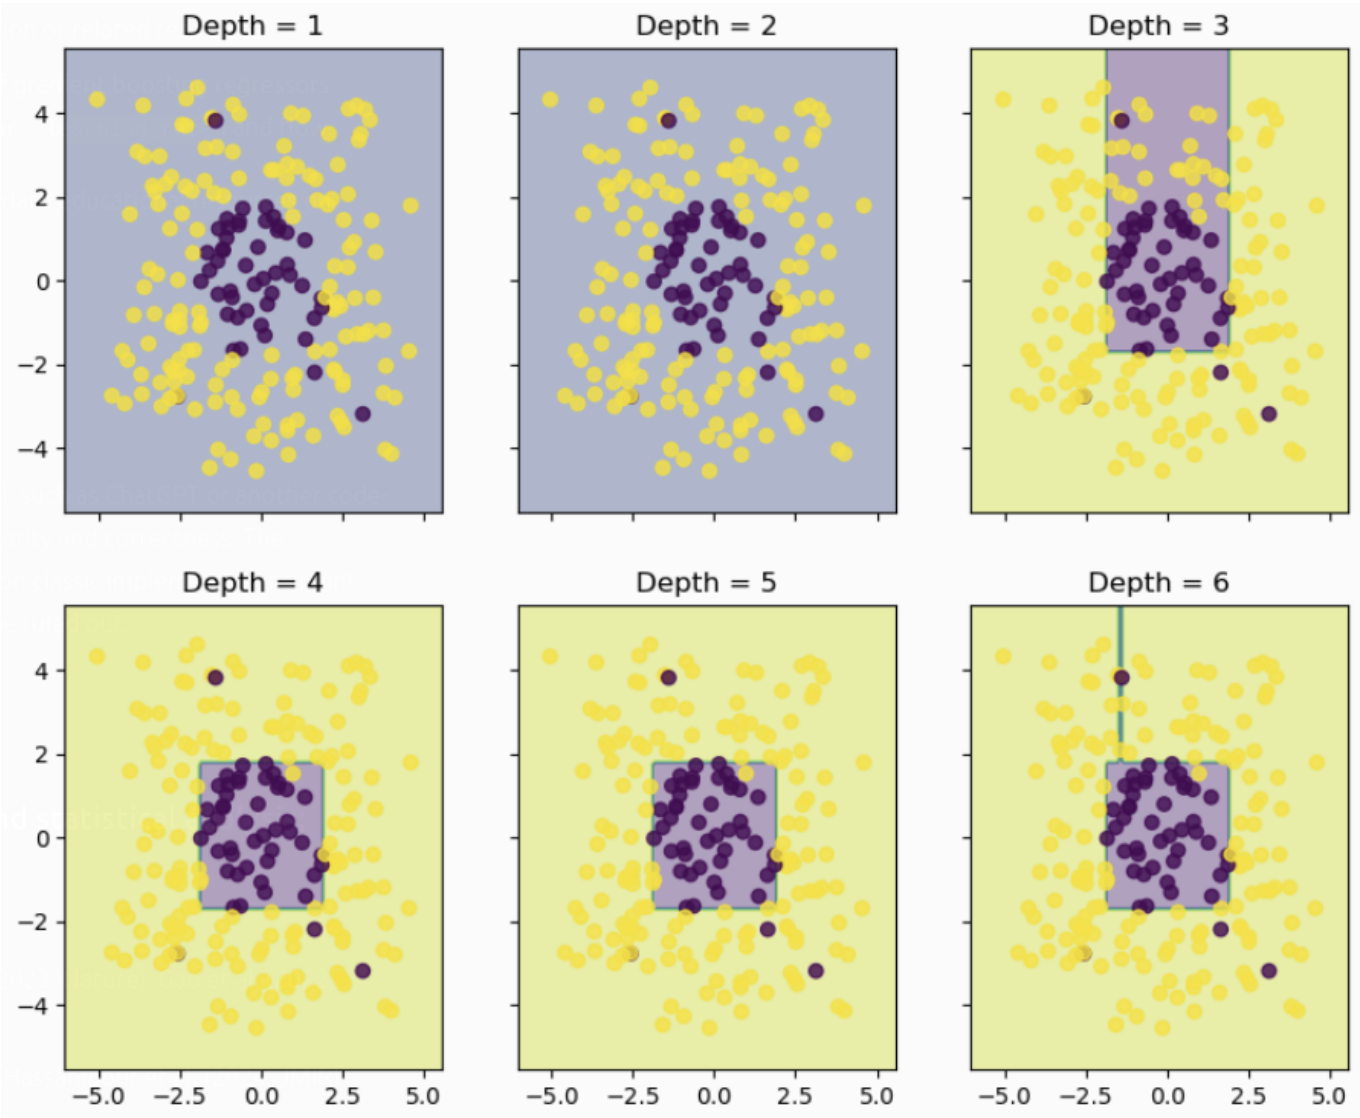
for idx, clf, tt in zip(product([0, 1], [0, 1, 2]),
                        [clf1, clf2, clf3, clf4, clf5, clf6],
                        ['Depth = {}'.format(n) for n in range(1, 7)]):

    Z = np.array([clf.predict_instance(x) for x in \
                  np.c_[xx.ravel(), yy.ravel()]])
    Z = Z.reshape(xx.shape)

    axarr[idx[0], idx[1]].contourf(xx, yy, Z, alpha=0.4)
    axarr[idx[0], idx[1]].scatter(x_train[:, 0], x_train[:, 1], \
                                  c=y_train_label[:, 0], alpha=0.8)
    axarr[idx[0], idx[1]].set_title(tt)

plt.show()

```



please turn over

Comparing decision tree with tree model sklearn:

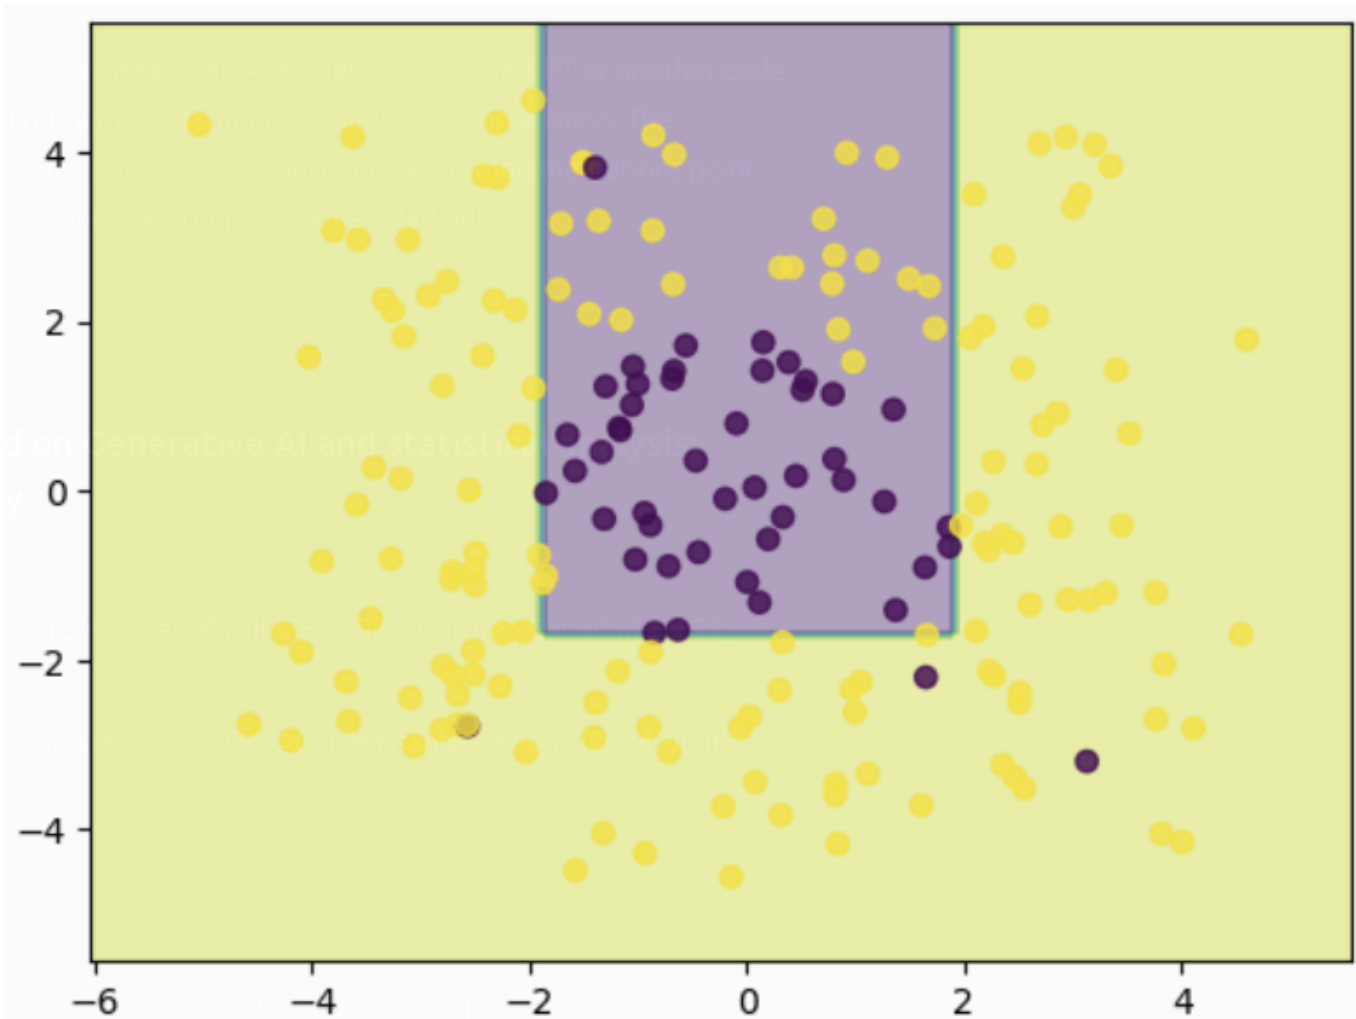
```

clf = DecisionTreeClassifier(criterion='entropy', max_depth=3, \
                             min_samples_split=5)
clf.fit(x_train, y_train_label)
export_graphviz(clf, out_file='tree_classifier.dot')

# Plotting decision regions
x_min, x_max = x_train[:, 0].min() - 1, x_train[:, 0].max() + 1
y_min, y_max = x_train[:, 1].min() - 1, x_train[:, 1].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.1),
                     np.arange(y_min, y_max, 0.1))

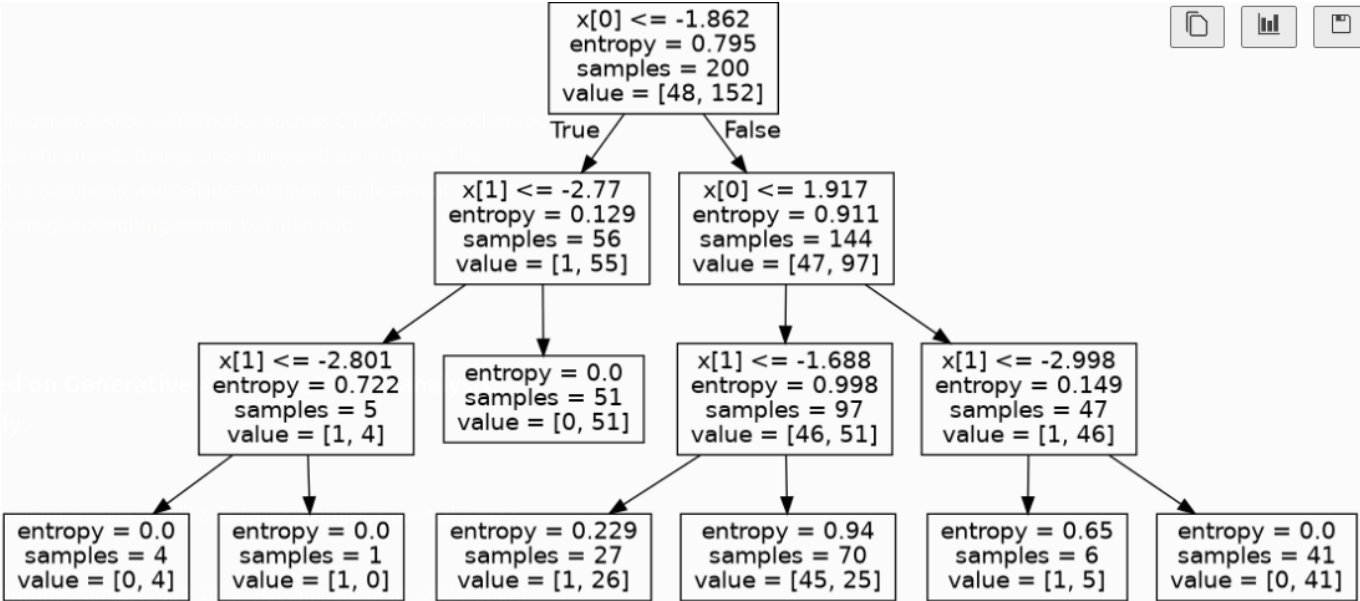
Z = np.array([clf.predict(x[np.newaxis,:]) \
              for x in np.c_[xx.ravel(), yy.ravel()]])
Z = Z.reshape(xx.shape)
plt.figure()
plt.contourf(xx, yy, Z, alpha=0.4)
plt.scatter(x_train[:, 0], x_train[:, 1],
            c=y_train_label[:,0], alpha=0.8)

```



please turn over

```
!dot -Tpng tree_classifier.dot -o tree_classifier.png
Image(filename='tree_classifier.png')
```



please turn over

Ans. 4

```
# Regression Tree Specific Code
def mean_absolute_deviation_around_median(y):
    '''
    Calculate the mean absolute deviation around the median of a \
    given target list

    :param y: a numpy array of targets shape = (n, 1)
    :return mae
    '''
    median = np.median(y)
    mae = np.mean(np.abs(y - median))

    return mae
```

Regression_Tree class and visualization code

```
class Regression_Tree():
    '''
    :attribute loss_function_dict: dictionary containing the loss \
    functions used for splitting
    :attribute estimator_dict: dictionary containing the estimation \
    functions used in leaf nodes
    '''

    loss_function_dict = {
        'mse': np.var,
        'mae': mean_absolute_deviation_around_median
    }

    estimator_dict = {
        'mean': np.mean,
        'median': np.median
    }

    def __init__(self, loss_function='mse', estimator='mean', \
        min_sample=5, max_depth=10):
        '''
        Initialize Regression_Tree
        :param loss_function(str): loss function used for splitting \
        internal nodes
        :param estimator(str): value estimator of internal node
        '''

        self.tree = Decision_Tree(\
            self.loss_function_dict[loss_function],
            self.estimator_dict[estimator],
            0, min_sample, max_depth)
```

```

def fit(self, X, y=None):
    self.tree.fit(X,y)
    return self

def predict_instance(self, instance):
    value = self.tree.predict_instance(instance)
    return value

data_krr_train = np.loadtxt('krr-train.txt')
data_krr_test = np.loadtxt('krr-test.txt')
x_krr_train, y_krr_train = data_krr_train[:,0].reshape(-1,1), \
    data_krr_train[:,1].reshape(-1,1)
x_krr_test, y_krr_test = data_krr_test[:,0].reshape(-1,1), \
    data_krr_test[:,1].reshape(-1,1)

# Training regression trees with different depth
clf1 = Regression_Tree(max_depth=1, min_sample=3, \
    loss_function='mae', estimator='mean')
clf1.fit(x_krr_train, y_krr_train)

clf2 = Regression_Tree(max_depth=2, min_sample=3, \
    loss_function='mae', estimator='mean')
clf2.fit(x_krr_train, y_krr_train)

clf3 = Regression_Tree(max_depth=3, min_sample=3, \
    loss_function='mae', estimator='mean')
clf3.fit(x_krr_train, y_krr_train)

clf4 = Regression_Tree(max_depth=4, min_sample=3, \
    loss_function='mae', estimator='mean')
clf4.fit(x_krr_train, y_krr_train)

clf5 = Regression_Tree(max_depth=5, min_sample=3, \
    loss_function='mae', estimator='mean')
clf5.fit(x_krr_train, y_krr_train)

clf6 = Regression_Tree(max_depth=10, min_sample=3, \
    loss_function='mae', estimator='mean')
clf6.fit(x_krr_train, y_krr_train)

plot_size = 0.001
x_range = np.arange(0., 1., plot_size).reshape(-1, 1)

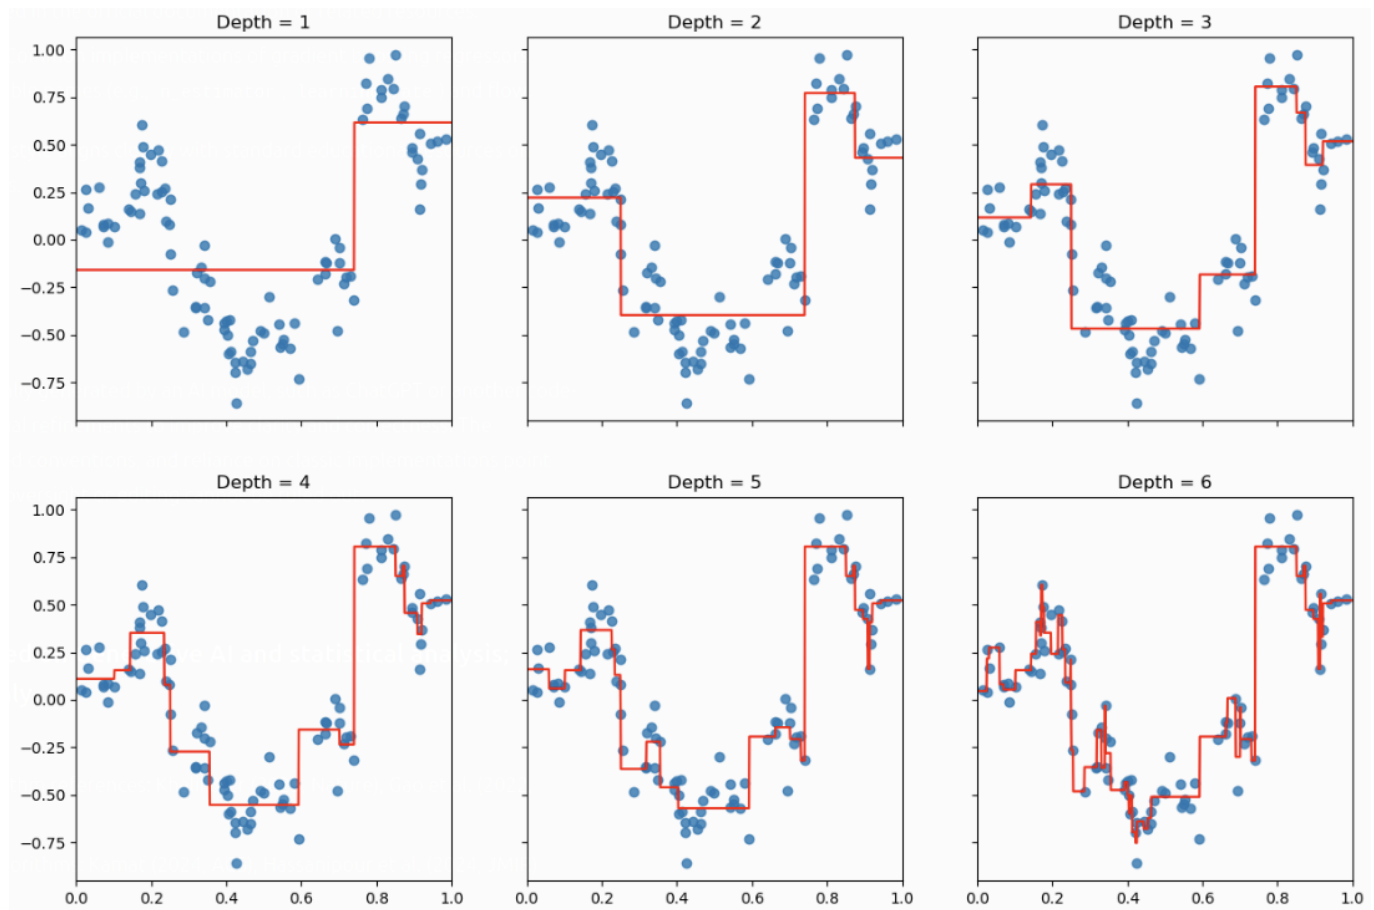
f2, axarr2 = plt.subplots(2, 3, sharex='col', \
    sharey='row', figsize=(15, 10))

for idx, clf, tt in zip(product([0, 1], [0, 1, 2]),
    [clf1, clf2, clf3, clf4, clf5, clf6],
    ['Depth = {}'.format(n) \
        for n in range(1, 7)]):

    y_range_predict = np.array([clf.predict_instance(x) \
        for x in x_range]).reshape(-1, 1)

```

```
axarr2[idx[0], idx[1]].plot(x_range, y_range_predict, color='r')
axarr2[idx[0], idx[1]].scatter(x_krr_train, y_krr_train, \
                               alpha=0.8)
axarr2[idx[0], idx[1]].set_title(tt)
axarr2[idx[0], idx[1]].set_xlim(0, 1)
plt.show()
```



please turn over

Ensembling

Ans 5.

```
#Pseudo-residual function.

def pseudo_residual_L2(train_target, train_predict):
    """
    Compute the pseudo-residual based on current predicted value.
    """
    return train_target - train_predict

class gradient_boosting():
    """
    Gradient Boosting regressor class
    :method fit: fitting model
    """
    def __init__(self, n_estimator, pseudo_residual_func,
learning_rate=0.01,
                min_sample=5, max_depth=5):
        """
        Initialize gradient boosting class

        :param n_estimator: number of estimators \
            (i.e. number of rounds of gradient boosting)
        :pseudo_residual_func: function used for computing pseudo-residual
between \
            training labels and predicted labels at each iteration
        :param learning_rate: step size of gradient descent
        """
        self.n_estimator = n_estimator
        self.pseudo_residual_func = pseudo_residual_func
        self.learning_rate = learning_rate
        self.min_sample = min_sample
        self.max_depth = max_depth

        self.estimators = [] #will collect the n_estimator models

    def fit(self, train_data, train_target):
        """
        Fit gradient boosting model
        :train_data array of inputs of size (n_samples, m_features)
        :train_target array of outputs of size (n_samples,)
        """
        self.base_model = DecisionTreeRegressor(criterion= 'squared_error',
\
            max_depth= self.max_depth, min_samples_leaf= self.min_sample)
        self.base_model.fit(train_data, train_target)

        for i in range(self.n_estimator):
            pred = self.learning_rate * self.base_model.predict(train_data)
```

```

        for model in self.estimators:
            pred = pred + self.learning_rate *
model.predict(train_data)

        residual = self.pseudo_residual_func(train_target, pred)

        h_m = DecisionTreeRegressor(max_depth= self.max_depth, \
            min_samples_leaf= self.min_sample)
        h_m.fit(train_data, residual)
        self.estimators.append(h_m)

def predict(self, test_data):
    """
    Predict value
    :train_data array of inputs of size (n_samples, m_features)
    """

    pred = self.base_model.predict(test_data)
    for model in self.estimators:
        pred = pred + (self.learning_rate * model.predict(test_data))
    return pred

```

Ans. 6

- (NOTE: I have built the gradient_boosting_class using sklearn DecisionTreeRegressor (code provided in the previous question), and using the Regression_Tree class (built upon the self-defined Decision_Tree class) (used in this question). This has been done, because it was unclear to me in Q.s 5 and 6 on what function to use for the gradient boosting models)

```

class gradient_boosting_v2():
    """
    Gradient Boosting regressor class
    :method fit: fitting model
    """
    def __init__(self, n_estimator, pseudo_residual_func,
learning_rate=0.01,
                min_sample=5, max_depth=5):
        """
        Initialize gradient boosting class

        :param n_estimator: number of estimators (i.e. number of rounds of
gradient boosting)
        :pseudo_residual_func: function used for computing pseudo-residual
between training labels and predicted labels at each iteration
        :param learning_rate: step size of gradient descent
        """
        self.n_estimator = n_estimator
        self.pseudo_residual_func = pseudo_residual_func
        self.learning_rate = learning_rate

```



```

        self.min_sample = min_sample
        self.max_depth = max_depth

        self.estimators = [] #will collect the n_estimator models

    def fit(self, train_data, train_target):
        """
        Fit gradient boosting model
        :train_data array of inputs of size (n_samples, m_features)
        :train_target array of outputs of size (n_samples,)
        """
        self.base_model = Regression_Tree(max_depth= self.max_depth,
min_sample= self.min_sample)
        self.base_model.fit(train_data, train_target)

        for i in range(self.n_estimator):
            pred = self.learning_rate *
np.array([self.base_model.predict_instance(instance) for instance in
train_data])
            for model in self.estimators:
                pred = pred + self.learning_rate *
np.array([model.predict_instance(instance) for instance in train_data])

            residual = self.pseudo_residual_func(train_target, pred)

            # h_m = DecisionTreeRegressor(max_depth= self.max_depth,
min_samples_leaf= self.min_sample)
            h_m = Regression_Tree(min_sample= self.min_sample, max_depth=
self.max_depth)
            h_m.fit(train_data, residual)
            self.estimators.append(h_m)

    def predict(self, test_data):
        """
        Predict value
        :train_data array of inputs of size (n_samples, m_features)
        """
        pred = np.array([self.base_model.predict_instance(instance) for
instance in test_data])
        for model in self.estimators:
            pred = pred + (self.learning_rate *
np.array([model.predict_instance(instance) for instance in test_data]))
        return pred

```

```

plot_size = 0.001
x_range = np.arange(0., 1., plot_size).reshape(-1, 1)

f2, axarr2 = plt.subplots(2, 3, sharex='col', sharey='row', figsize=(15,
10))

```

```

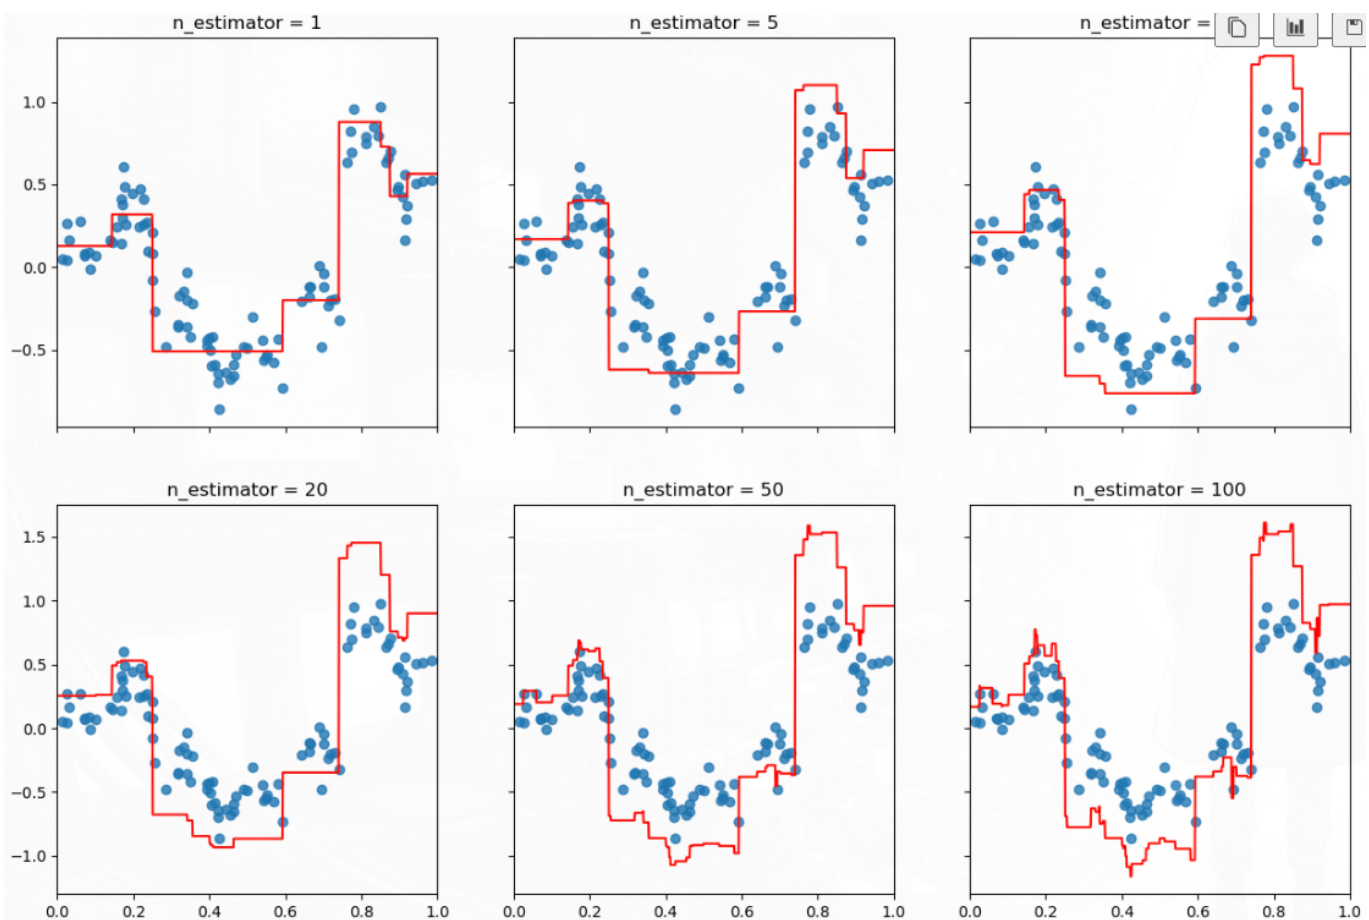
for idx, i, tt in zip(product([0, 1], [0, 1, 2]),
                        [1, 5, 10, 20, 50, 100],
                        ['n_estimator = {}'.format(n) \
                         for n in [1, 5, 10, 20, 50, 100]]):

    gbm_1d = gradient_boosting_v2(n_estimator=i,
                                  pseudo_residual_func=pseudo_residual_L2, \
                                  max_depth=3, learning_rate=0.1)
    gbm_1d.fit(x_krr_train, y_krr_train[:,0])

    y_range_predict = gbm_1d.predict(x_range)

    axarr2[idx[0], idx[1]].plot(x_range, y_range_predict, color='r')
    axarr2[idx[0], idx[1]].scatter(x_krr_train, y_krr_train, alpha=0.8)
    axarr2[idx[0], idx[1]].set_title(tt)
    axarr2[idx[0], idx[1]].set_xlim(0, 1)

```



debugging code via sklearn

```

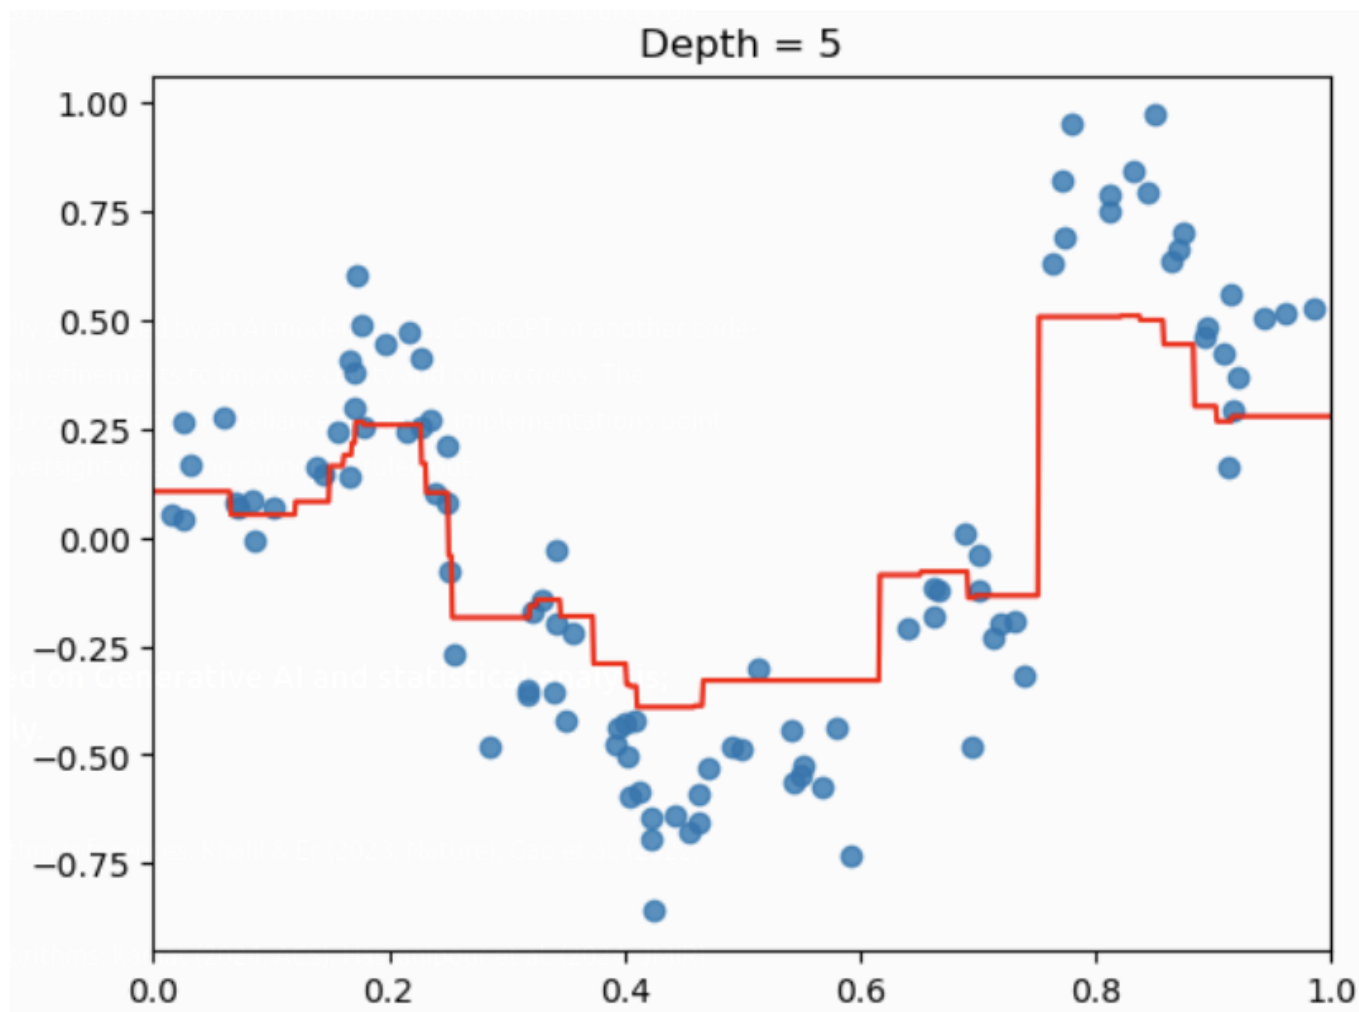
plot_size = 0.001
x_range = np.arange(0., 1., plot_size).reshape(-1, 1)
max_depth = 5

clf = GradientBoostingRegressor(criterion='squared_error', \
                                max_depth=max_depth, min_samples_leaf=5, learning_rate= 0.01,

```

```
n_estimators= 100)
clf.fit(x_krr_train, y_krr_train)
y_range_predict = np.array([clf.predict(x[np.newaxis,:]) for x in
x_range]).reshape(-1, 1)

plt.plot(x_range, y_range_predict, color='r')
plt.scatter(x_krr_train, y_krr_train, alpha=0.8)
plt.title('Depth = {}'.format(max_depth))
plt.xlim(0, 1)
plt.show()
```



please turn over

Neural Network Introduction, Computation Graph Framework

- No questions asked here please turn over

Ridge Regression

Ans. 7

```
class L2NormPenaltyNode(object):
    """ Node computing  $l2\_reg * ||w||^2$  for scalars  $l2\_reg$  and vector  $w$  """
    def __init__(self, l2_reg, w, node_name):
        """
        Parameters:
        l2_reg: a numpy scalar array (e.g. np.array(.01)) (not a node)
        w: a node for which w.out is a numpy vector
        node_name: node's name (a string)
        """
        self.node_name = node_name
        self.out = None
        self.d_out = None
        self.l2_reg = np.array(l2_reg)
        self.w = w

    def forward(self):
        w_vec = self.w.out
        w_norm = np.linalg.norm(w_vec)
        self.out = ((w_norm * w_norm) * self.l2_reg)
        self.d_out = np.zeros_like(self.out)
        return self.out

    def backward(self):
        w_vec = self.w.out
        self.w.d_out = 2 * self.l2_reg * w_vec * self.d_out

    def get_predecessors(self):
        return [self.w]
```

```
• (deep_learning) (base) arjun-prasad@arjun-prasad-Lenovo-Legion-5-15ARH05:~/ARJUN/CourseWork/Mach
DEBUG: (Node l2 norm node) Max rel error for partial deriv w.r.t. w is 4.190667061846293e-09.
```

please turn over

Ans 8.

```
class SumNode(object):
    """ Node computing a + b, for numpy arrays a and b"""
    def __init__(self, a, b, node_name):
        """
        Parameters:
        a: node for which a.out is a numpy array
        b: node for which b.out is a numpy array of the same shape as a
        node_name: node's name (a string)
        """
        self.node_name = node_name
        self.out = None
        self.d_out = None
        self.b = b
        self.a = a

    def forward(self):
        self.out = self.a.out + self.b.out
        self.d_out = np.zeros_like(self.out)
        return self.out

    def backward(self):
        d_a = 1 * self.d_out
        d_b = 1 * self.d_out
        self.a.d_out += d_a
        self.b.d_out += d_b

    def get_predecessors(self):
        return [self.a, self.b]
```

```
.DEBUG: (Node sum node) Max rel error for partial deriv w.r.t. a is 1.636578803425771e-09.
DEBUG: (Node sum node) Max rel error for partial deriv w.r.t. b is 5.838672347430118e-10.
```

please turn over

Ans 9.

```

class RidgeRegression(BaseEstimator, RegressorMixin):
    """ Ridge regression with computation graph """
    def __init__(self, l2_reg=1, step_size=.005, max_num_epochs = 5000):
        self.max_num_epochs = max_num_epochs
        self.step_size = step_size

        # Build computation graph
        self.x = nodes.ValueNode(node_name="x") # to hold a vector input
        self.y = nodes.ValueNode(node_name="y") # to hold a scalar response
        self.w = nodes.ValueNode(node_name="w") # to hold the parameter
        vector
        self.b = nodes.ValueNode(node_name="b") # to hold the bias
        parameter (scalar)
        self.prediction = nodes.VectorScalarAffineNode(\
            x=self.x, \
            w=self.w, b=self.b, \
            node_name="prediction")
        self.objective = nodes.SumNode(nodes.SquaredL2DistanceNode(\
            a=self.prediction, b=self.y, \
            node_name="square loss"), \
            nodes.L2NormPenaltyNode(l2_reg, self.w, "l2_reg"), \
            "ridge_reg")

        self.inputs = [self.x]
        self.outcomes = [self.y]
        self.parameters = [self.w, self.b]

        self.graph = graph.ComputationGraphFunction(self.inputs, \
            self.outcomes, self.parameters, \
            self.prediction, self.objective)

    def fit(self, X, y):
        num_instances, num_ftrs = X.shape
        y = y.reshape(-1)

        init_parameter_values = {"w": np.zeros(num_ftrs), "b":
np.array(0.0)}
        self.graph.set_parameters(init_parameter_values)

        for epoch in range(self.max_num_epochs):
            shuffle = np.random.permutation(num_instances)
            epoch_obj_tot = 0.0
            for j in shuffle:
                obj, grads = self.graph.get_gradients(input_values = {"x":
X[j]},
                                                    outcome_values = {"y":
y[j]})
                #print(obj)
                epoch_obj_tot += obj

```

```

        # Take step in negative gradient direction
        steps = {}
        for param_name in grads:
            steps[param_name] = -self.step_size * grads[param_name]
        self.graph.increment_parameters(steps)

    if epoch % 50 == 0:
        train_loss = sum((y - self.predict(X,y)) **2)/num_instances
        print("Epoch ",epoch," : Ave objective=",\
              epoch_obj_tot/num_instances,\
              " Ave training loss: ",train_loss)

def predict(self, X, y=None):
    try:
        getattr(self, "graph")
    except AttributeError:
        raise RuntimeError("You must train classifier before predicting
data!")

    num_instances = X.shape[0]
    preds = np.zeros(num_instances)
    for j in range(num_instances):
        preds[j] = self.graph.get_prediction(input_values={"x":X[j]})

    return preds

def main():
    data_fname = "data.pickle"
    x_train, y_train, x_val, y_val, target_fn, coefs_true, \
        featurize = setup_problem.load_problem(data_fname)

    # Generate features
    X_train = featurize(x_train)
    X_val = featurize(x_val)

    pred_fns = []
    x = np.sort(np.concatenate([np.arange(0,1,.001), x_train]))
    X = featurize(x)

    l2reg = 1
    estimator = RidgeRegression(l2_reg=l2reg, \
        step_size=0.00005, max_num_epochs=2000)
    estimator.fit(X_train, y_train)
    name = "Ridge with L2Reg="+str(l2reg)
    pred_fns.append({"name":name, "preds": estimator.predict(X) })

    l2reg = 0
    estimator = RidgeRegression(l2_reg=l2reg, \
        step_size=0.0005, max_num_epochs=500)
    estimator.fit(X_train, y_train)
    name = "Ridge with L2Reg="+str(l2reg)

```



```

    pred_fns.append({"name": name, "preds": estimator.predict(X) })

    # Let's plot prediction functions and compare coefficients for several
fits
    # and the target function.

    pred_fns.append({"name": "Target Parameter Values (i.e. Bayes
Optimal)", \
        "coefs": coefs_true, "preds": target_fn(x)})

    plot_utils.plot_prediction_functions(x, pred_fns, x_train, \
        y_train, legend_loc="best")

if __name__ == '__main__':
    main()

```

- test cases

```

• (deep_learning) (base) arjun-prasad@arjun-prasad-Lenovo-Legion-5-15ARH05:~/ARJUN/Cour
idge_regression.t.py
DEBUG: (Node l2 norm node) Max rel error for partial deriv w.r.t. w is 5.628854579506
.DEBUG: (Node sum node) Max rel error for partial deriv w.r.t. a is 5.838671770815658
DEBUG: (Node sum node) Max rel error for partial deriv w.r.t. b is 1.6365788599345556
.DEBUG: (Parameter w) Max rel error for partial deriv 2.7820614265556677e-09.
DEBUG: (Parameter b) Max rel error for partial deriv 7.940832693784278e-10.

```

- Average Training Loss:

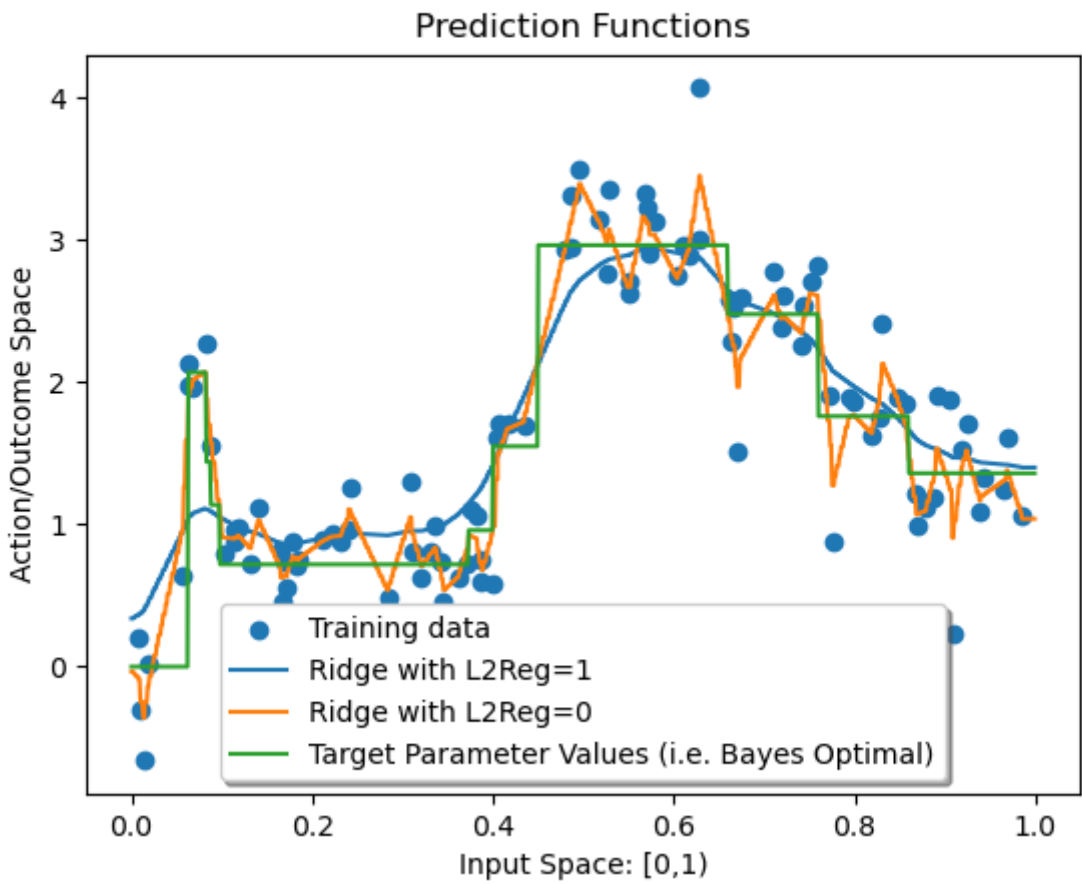
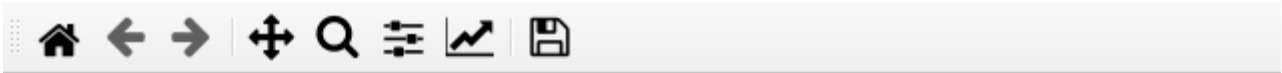
```

Epoch 1750 : Ave objective= 0.30327767702891556 Ave training loss: 0.1995312718970577
Epoch 1800 : Ave objective= 0.30422838177646816 Ave training loss: 0.19951172464135844
Epoch 1850 : Ave objective= 0.3051244963548033 Ave training loss: 0.19948279848411762
Epoch 1900 : Ave objective= 0.30458070045385155 Ave training loss: 0.20049564456872507
Epoch 1950 : Ave objective= 0.3049240603230615 Ave training loss: 0.19998102561296593
Epoch 0 : Ave objective= 0.6909674276953551 Ave training loss: 0.5351599062284119
Epoch 50 : Ave objective= 0.11900738345274707 Ave training loss: 0.10623783437136644
Epoch 100 : Ave objective= 0.10561873811273172 Ave training loss: 0.08500256731032349
Epoch 150 : Ave objective= 0.08378111519194606 Ave training loss: 0.07946335016874015
Epoch 200 : Ave objective= 0.0718557501028382 Ave training loss: 0.06497773902021574
Epoch 250 : Ave objective= 0.07118748491708027 Ave training loss: 0.06323625558006388
Epoch 300 : Ave objective= 0.05812454069791321 Ave training loss: 0.05272421376954614
Epoch 350 : Ave objective= 0.05478289328024275 Ave training loss: 0.05396470524099741
Epoch 400 : Ave objective= 0.05430683088238369 Ave training loss: 0.050801303495726716
Epoch 450 : Ave objective= 0.050543836008525085 Ave training loss: 0.05287872115369579

```

- Avg training loss 1, l2reg = 1 => 0.19998102561296593
- Avg training loss 2, l2reg = 0 => 0.05287872115369579

- output image generated on running `python ridge_regression.py`



please turn over

Multilayer Perceptron

Ans 10.

(10) To prove: $\frac{\partial J}{\partial w_{ij}^0} = \frac{\partial J}{\partial y_i^0} x_j^0$ $\parallel x = (x_1, x_2, \dots, x_d)^T$

Proof:

for 'n' samples

$$\frac{\partial J}{\partial w_{ij}^0} = \sum_{k=1}^n \frac{\partial J}{\partial y_k} \cdot \frac{\partial y_k}{\partial w_{ij}^0}$$

$$= 0 + \underbrace{\sum_{k \neq i} \frac{\partial J}{\partial y_k} \left(\frac{\partial y_k}{\partial w_{ij}^0} \right)}_{\text{for } k \neq i} \quad (\forall k \in \{1, 2, \dots, n\})$$

$$\rightarrow \frac{\partial y_i^0}{\partial w_{ij}^0} = \frac{\partial}{\partial w_{ij}^0} \left(\sum_{j=1}^d w_{ij}^0 x_j^0 + b_i^0 \right)$$

$$= x_j^0$$

Therefore

$$\boxed{\frac{\partial J}{\partial w_{ij}^0} = 0 + \frac{\partial J}{\partial y_i^0} x_j^0}$$

∴ Hence Proved.

please turn over

Ans 11.

$$\textcircled{11} \quad \frac{\partial J}{\partial y} \in \mathbb{R}^{m \times 1}, \text{ row entry} \Rightarrow \frac{\partial J}{\partial y_i}$$

$$\frac{\partial J}{\partial W} \in \mathbb{R}^{m \times d}, \text{ } i, j \text{th entry} \Rightarrow \frac{\partial J}{\partial w_{ij}}$$

We know that for a given sample x ,

$$\frac{\partial J}{\partial w_{ij}} = \frac{\partial J}{\partial y_i} \cdot x_j \quad \left\{ \begin{array}{l} \text{(from prob. 10)} \end{array} \right.$$

\therefore for $X \in \mathbb{R}^{m \times d}$ (m samples, d features)

we can write

$$\frac{\partial J}{\partial w_{ij}} = \frac{\partial J}{\partial y_i} \cdot x_{ij}$$

Thus

$$\underbrace{\left(\frac{\partial J}{\partial W} \right)^T}_{d \times m} = \underbrace{X}_{m \times d} \cdot \underbrace{\left(\frac{\partial J}{\partial y} \right)^T}_{m \times 1}$$

$$\Rightarrow \boxed{\frac{\partial J}{\partial W} = \frac{\partial J}{\partial y} \cdot X^T}$$

(outer product)

$$\left[(AB)^T = B^T A^T \right]$$

please turn over

Ans 12.

(12)

$$\frac{\partial J}{\partial x} \in \mathbb{R}^d$$

$$\frac{\partial J}{\partial x} = \begin{bmatrix} \frac{\partial J}{\partial x_1} \\ \frac{\partial J}{\partial x_2} \\ \vdots \\ \frac{\partial J}{\partial x_d} \end{bmatrix}$$

$$\frac{\partial J}{\partial x_k} = \sum_{i=1}^m \frac{\partial J}{\partial y_i^o} \frac{\partial y_i^o}{\partial x_k}.$$

$$y_i^o = \sum_{l=1}^d w_{il}^o x_l + b_i^o$$

$$\Rightarrow \frac{\partial y_i^o}{\partial x_k} = \sum_{l=1}^d \frac{\partial (w_{il}^o x_l)}{\partial x_k} + b_i^o$$

$$\boxed{\frac{\partial y_i^o}{\partial x_k} = w_{ik}^o.}$$

$$\therefore \frac{\partial J}{\partial x_k} = \sum_{i=1}^m \frac{\partial J}{\partial y_i^o} \cdot w_{ik}^o.$$

$$\Rightarrow \boxed{\frac{\partial J}{\partial x} = W^T \frac{\partial J}{\partial y}} \quad \text{+ hence proved}$$

please turn over

Ans 13.

(13)

$$y = wx + b$$
$$\Rightarrow \frac{\partial y}{\partial b} = 1.$$
$$\frac{\partial J}{\partial b} = \frac{\partial J}{\partial y} \cdot \frac{\partial y}{\partial b} \rightarrow 1$$

$\frac{\partial J}{\partial b} = \frac{\partial J}{\partial y}$

 hence proved.

please turn over

Ans 14.

(14)

$$S = \sigma(A)$$

for i^{th} dimension of A , we have A_i^0 .

$\therefore \downarrow$

$$\frac{\partial J}{\partial A_i^0} = \frac{\partial J}{\partial (\sigma(A_i^0))} \cdot \frac{\partial (\sigma(A_i^0))}{\partial A_i^0}$$

$$= \frac{\partial J}{\partial s_i^0} \cdot \sigma'(A_i^0)$$

\Rightarrow

$$\boxed{\frac{\partial J}{\partial A} = \frac{\partial J}{\partial s_i^0} \odot \sigma'(A_i^0)}$$

please turn over

Ans 15.

```
class AffineNode(object):
    """Node implementing affine transformation (W,x,b)-->Wx+b, where W is a
    matrix,
    and x and b are vectors
    Parameters:
    W: node for which W.out is a numpy array of shape (m,d)
    x: node for which x.out is a numpy array of shape (d)
    b: node for which b.out is a numpy array of shape (m) \
        (i.e. vector of length m)
    """
    def __init__(self, w, x, b, node_name):
        self.node_name = node_name
        self.out = None
        self.d_out = None
        self.x = x
        self.w = w
        self.b = b

    def forward(self):
        self.out = np.dot(self.w.out, self.x.out) + self.b.out
        self.d_out = np.zeros(self.out.shape)
        return self.out

    def backward(self):
        d_x = np.dot(self.w.out.T, self.d_out)
        d_b = self.d_out
        d_w = np.outer(self.d_out, self.x.out)
        self.x.d_out += d_x
        self.w.d_out += d_w
        self.b.d_out += d_b

    def get_predecessors(self):
        return [self.x, self.w, self.b]
```

test results

```
DEBUG: (Node affine) Max rel error for partial deriv w.r.t. x is 5.298097672432879e-09.
DEBUG: (Node affine) Max rel error for partial deriv w.r.t. W is 1.694796002711163e-09.
DEBUG: (Node affine) Max rel error for partial deriv w.r.t. b is 1.6365789068091537e-09.
```

please turn over

Ans. 16

```
class TanhNode(object):
    """Node tanh(a), where tanh is applied elementwise \
    to the array a
    Parameters:
    a: node for which a.out is a numpy array
    """
    def __init__(self, a, node_name):
        self.node_name = node_name
        self.out = None
        self.d_out = None
        self.a = a

    def forward(self):
        self.out = np.tanh(self.a.out)
        self.d_out = np.zeros(self.out.shape)
        return self.out

    def backward(self):
        d_a = (1-self.out**2)*self.d_out
        self.a.d_out += d_a
        return self.d_out

    def get_predecessors(self):
        return [self.a]
```

- test results

```
.DEBUG: (Node tanh) Max rel error for partial deriv w.r.t. a is 3.647110392742941e-08.
```

please turn over

Ans. 17

```

class MLPRegression(BaseEstimator, RegressorMixin):
    """ MLP regression with computation graph """
    def __init__(self, num_hidden_units=10, step_size=.005, \
        init_param_scale=0.01, max_num_epochs = 5000):
        self.num_hidden_units = num_hidden_units
        self.init_param_scale = init_param_scale
        self.max_num_epochs = max_num_epochs
        self.step_size = step_size

        self.x = nodes.ValueNode(node_name="x") # to hold a vector input
        self.y = nodes.ValueNode(node_name="y") # to hold a scalar response
        self.w1 = nodes.ValueNode(node_name="w1") # to hold the parameter
vector
        self.b1 = nodes.ValueNode(node_name=f"b1") # to hold the bias
parameter (scalar)

        self.w2 = nodes.ValueNode(node_name="w2") # to hold the parameter
vector
        self.b2 = nodes.ValueNode(node_name=f"b2")

        self.prediction = nodes.AffineNode(x = self.x, w = self.w1, \
            b = self.b1, node_name="hidden_layer")
        self.prediction = nodes.TanhNode(a = self.prediction,
node_name=f"activation")

        self.prediction = nodes.VectorScalarAffineNode(x = self.prediction
, \
            w=self.w2, b = self.b2, node_name="output_layer")

        self.objective = nodes.SquaredL2DistanceNode(a=self.prediction, \
            b=self.y, node_name="square loss")

        # Group nodes into types to construct computation graph function
        self.inputs = [self.x]
        self.outcomes = [self.y]
        self.parameters = [self.w1, self.w2, self.b1, self.b2]

        self.graph = graph.ComputationGraphFunction(self.inputs, \
            self.outcomes, self.parameters, self.prediction, self.objective)

    def fit(self, X, y):
        num_instances, num_ftrs = X.shape
        y = y.reshape(-1)
        s = self.init_param_scale
        init_values = {"w1": s *
np.random.standard_normal((self.num_hidden_units, \
            num_ftrs)),
            "b1": s *
np.random.standard_normal((self.num_hidden_units)),

```

```

        "w2": s *
np.random.standard_normal((self.num_hidden_units)),
        "b2": s * np.array(np.random.randn()) }

self.graph.set_parameters(init_values)

for epoch in range(self.max_num_epochs):
    shuffle = np.random.permutation(num_instances)
    epoch_obj_tot = 0.0
    for j in shuffle:
        obj, grads = self.graph.get_gradients(input_values = {"x":
X[j]},
                                                outcome_values = {"y":
y[j]})

        #print(obj)
        epoch_obj_tot += obj
        # Take step in negative gradient direction
        steps = {}
        for param_name in grads:
            steps[param_name] = -self.step_size * grads[param_name]
        self.graph.increment_parameters(steps)
        #pdb.set_trace()

    if epoch % 50 == 0:
        train_loss = sum((y - self.predict(X,y)) **2)/num_instances
        print("Epoch ",epoch,": Ave objective=",\
              epoch_obj_tot/num_instances," Ave training loss:
",train_loss)

    def predict(self, X, y=None):
        try:
            getattr(self, "graph")
        except AttributeError:
            raise RuntimeError("You must train classifier before predicting
data!")

        num_instances = X.shape[0]
        preds = np.zeros(num_instances)
        for j in range(num_instances):
            preds[j] = self.graph.get_prediction(input_values={"x":X[j]})

        return preds

```

- test cases

```
DEBUG: (Node affine) Max rel error for partial deriv w.r.t. x is 5.298097672432879e-09.
DEBUG: (Node affine) Max rel error for partial deriv w.r.t. W is 1.694796002711163e-09.
DEBUG: (Node affine) Max rel error for partial deriv w.r.t. b is 1.6365789068091537e-09.
.DEBUG: (Node tanh) Max rel error for partial deriv w.r.t. a is 3.647110392742941e-08.
.DEBUG: (Parameter W1) Max rel error for partial deriv 1.7553991632627748e-06.
DEBUG: (Parameter b1) Max rel error for partial deriv 2.1649821289952703e-07.
DEBUG: (Parameter w2) Max rel error for partial deriv 1.464398757659339e-09.
DEBUG: (Parameter b2) Max rel error for partial deriv 5.042356208756589e-10.
.
-----
Ran 3 tests in 0.003s

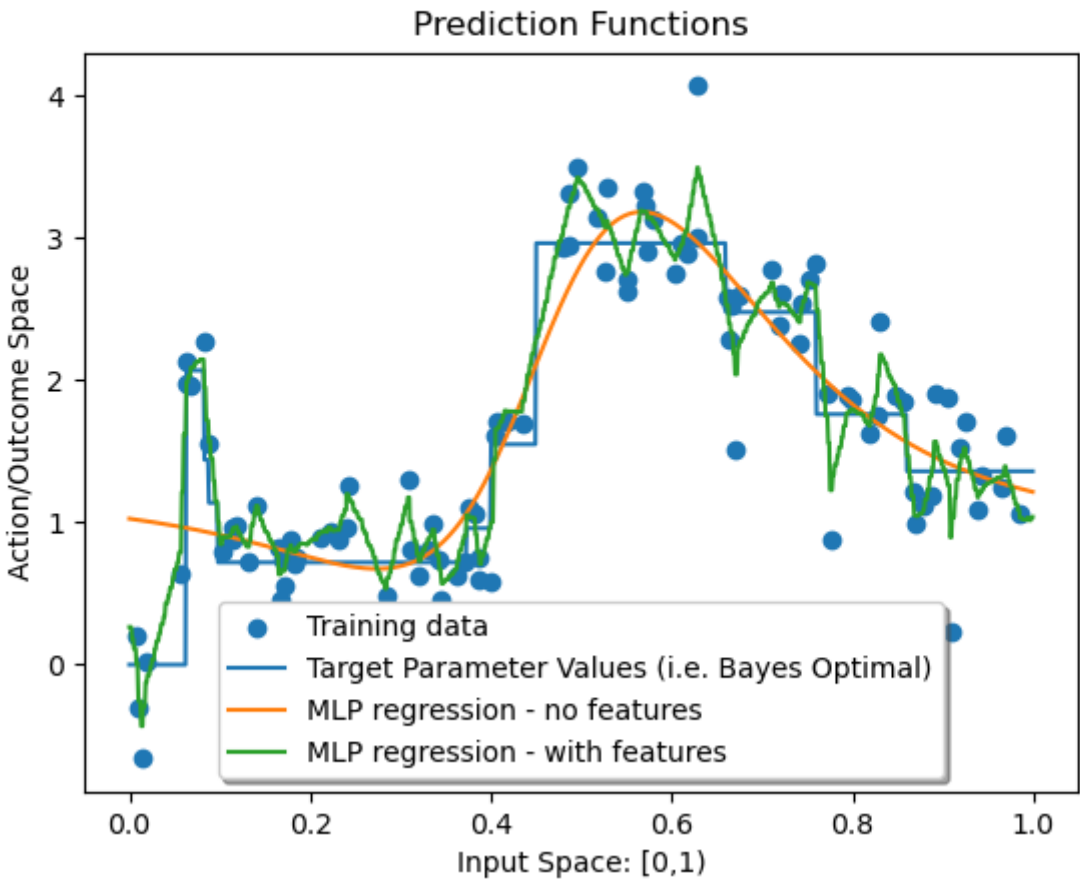
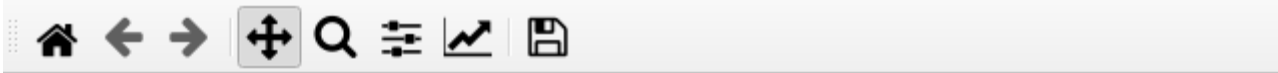
OK
```

- running `python mlp_regression.py`

```
Epoch 4800 : Ave objective= 0.24417850889609855 Ave training loss: 0.23907543751001983
Epoch 4850 : Ave objective= 0.24038347369486165 Ave training loss: 0.24084079771282443
Epoch 4900 : Ave objective= 0.24243009939525678 Ave training loss: 0.23829217509160958
Epoch 4950 : Ave objective= 0.24175358009397616 Ave training loss: 0.2375514899589099
Epoch 0 : Ave objective= 3.163273548673171 Ave training loss: 2.5605331414531265
Epoch 50 : Ave objective= 0.14844708437599397 Ave training loss: 0.1397147611961311
Epoch 100 : Ave objective= 0.1116268144045757 Ave training loss: 0.1223455989479196
Epoch 150 : Ave objective= 0.1025335662192338 Ave training loss: 0.09111530930261971
Epoch 200 : Ave objective= 0.08920208778374379 Ave training loss: 0.08274506292670711
Epoch 250 : Ave objective= 0.0770484544738569 Ave training loss: 0.06741359647846575
Epoch 300 : Ave objective= 0.06939536180896462 Ave training loss: 0.06978848314951241
Epoch 350 : Ave objective= 0.05688447211810779 Ave training loss: 0.07793590268056473
Epoch 400 : Ave objective= 0.056197564706299934 Ave training loss: 0.04543237206017478
Epoch 450 : Ave objective= 0.05079795148881582 Ave training loss: 0.04350373562087348
```

- avg training error 1 = 0.2375514899589099
- avg training error 2 = 0.04350373562087348

- output generated on running `python mlp_regression.py`



please turn over

- *NOTE: optional questions 18, 19, and 20 have not been added to this pdf*